

# kernel



# things programs on portal shouldn't do

read other user's files

modify OS's memory

read other user's data in memory

hang the entire system



# things programs on portal shouldn't do

*read other user's files*

modify OS's memory

read other user's data in memory

hang the entire system



# things programs on portal shouldn't do

read other user's files

*modify OS's memory*

*read other user's data in memory*

hang the entire system



# things programs on portal shouldn't do

read other user's files

modify OS's memory

read other user's data in memory

*hang the entire system*



# againframe(portalShouldnt)



# privileged operation: problem

how can hardware (HW) plus operating system (OS) allow:

read your own files from hard drive

but disallow:

read others files from hard drive



# some ideas

OS tells HW 'okay' parts of hard drive before running program code

complex for hardware and for OS

OS verifies your program's code can't do bad hard drive access

no work for HW, but complex for OS

may require compiling differently to allow analysis

OS tells HW to only allow OS-written code to access hard drive

that code can enforce only 'good' accesses

requires program code to call OS routines to access hard drive

relatively simple for hardware



# some ideas

OS tells HW 'okay' parts of hard drive before running program code

complex for hardware and for OS

OS verifies your program's code can't do bad hard drive access

no work for HW, but complex for OS

may require compiling differently to allow analysis

*OS tells HW to only allow OS-written code to access hard drive*

that code can enforce only 'good' accesses

requires program code to call OS routines to access hard drive

relatively simple for hardware



# kernel mode

extra one-bit register: “are we in kernel mode”

other names: privileged mode, supervisor mode, ...

not in kernel mode = user mode

certain operations only allowed in kernel mode

privileged instructions

example: talking to any I/O device



# what runs in kernel mode?

system boots in kernel mode


OS switches to user mode to run program code

next topic: when does system switch back to kernel mode?


how does OS tell HW where the (trusted) OS code is?



# hardware + system call interface




# hardware + system call + library interface



# againframe(unixDesignHWNoLib)



# calling the OS?



# controlled entry to kernel mode (1)

special instruction: “make system call”

similar idea as `call` instruction — jump to function elsewhere  
(and allow that function to return later)

runs OS code in kernel mode at location specified earlier

OS sets up at boot

location can't be changed without privileged instruction



# controlled entry to kernel mode (2)

OS needs to make specified location:

figure out what operation the program wants

calling convention, similar to function arguments + return value

be “safe” — not allow the program to do ‘bad’ things

example: checks whether current program is allowed to read file  
before reading it


requires exceptional care — program can try weird things




# system call process



# system call process



# system call process



# system call terminology

some inconsistency:

system call = event of entering kernel mode on request?

system call = whole process from beginning to end?

same issue as with 'function call'

is it just starting the function, or the whole time the function runs?



# Linux x86-64 system calls

special instruction: `syscall`

runs OS specified code in kernel mode



# Linux syscall calling convention

before `syscall`:

`%rax` — system call number

`%rdi, %rsi, %rdx, %r10, %r8, %r9` — args

after `syscall`:

`%rax` — return value

on error: `%rax` contains -1 times “error number”

*almost* the same as normal function calls



# Linux x86-64 hello world

```
.globl _start
.data
hello_str: .asciz "Hello, World!\n"
.text
_start:
    movq $1, %rax # 1 = "write"
    movq $1, %rdi # file descriptor 1 = stdout
    movq $hello_str, %rsi
    movq $15, %rdx # 15 = strlen("Hello, World!\n")
    syscall

    movq $60, %rax # 60 = exit
    movq $0, %rdi
    syscall
```



# approx. system call handler

```
sys_call_table:  
    .quad handle_read_syscall  
    .quad handle_write_syscall  
    // ...  
  
handle_syscall:  
    ... // save old PC, etc.  
    pushq %rcx // save registers  
    pushq %rdi  
    ...  
    call *sys_call_table(,%rax,8)  
    ...  
    popq %rdi  
    popq %rcx  
    return_from_exception
```



# Linux system call examples

`mmap`, `brk` — allocate memory

`fork` — create new process

`execve` — run a program in the current process

`open`, `read`, `write` — access files

`_exit` — terminate a process

`socket`, `accept`, `getpeername` — socket-related



# Linux system call examples

mmap, brk — allocate memory

fork — create new process

execve — run a program in the current process

open, *read*, write — access files

*\_exit* — terminate a process


socket, accept, getpeername — socket-related




# system call handled slowly?




# system call handled slowly?




# system call handled slowly?



# system call handled slowly?



# system call handled slowly?



# system call wrappers

library functions to not write assembly:

open:

```
movq $2, %rax // 2 = sys_open
// 2 arguments happen to use same registers
syscall
// return value in %eax
cmp $0, %rax
jl has_error
ret
```

has\_error:

```
neg %rax
movq %rax, errno
movq $-1, %rax
ret
```



# system call wrappers

library functions to not write assembly:

open:

```
movq $2, %rax // 2 = sys_open
// 2 arguments happen to use same registers
```

*syscall*

```
// return value in %eax
```

```
cmp $0, %rax
```

```
jnl has_error
```

```
ret
```

has\_error:

```
neg %rax
```

```
movq %rax, errno
```

```
movq $-1, %rax
```

```
ret
```



# system call wrappers

library functions to not write assembly:

open:

```
movq $2, %rax // 2 = sys_open
// 2 arguments happen to use same registers
syscall
// return value in %eax
cmp $0, %rax
jl has_error
ret
```

has\_error:

```
neg %rax
movq %rax, errno
movq $-1, %rax
ret
```



# system call wrapper: usage

```
/* unistd.h contains definitions of:
   O_RDONLY (integer constant), open() */
#include <unistd.h>
int main(void) {
    int file_descriptor;
    file_descriptor = open("input.txt", O_RDONLY);
    if (file_descriptor < 0) {
        printf("error: %s\\n", strerror(errno));
        exit(1);
    }
    ...
    result = read(file_descriptor, ...);
    ...
}
```



# system call wrapper: usage

```
/* unistd.h contains definitions of:
   O_RDONLY (integer constant), open() */
#include <unistd.h>
int main(void) {
    int file_descriptor;
    file_descriptor = open("input.txt", O_RDONLY);
    if (file_descriptor < 0) {
        printf("error: %s\\n", strerror(errno));
        exit(1);
    }
    ...
    result = read(file_descriptor, ...);
    ...
}
```



# system call wrapper: usage

```
/* unistd.h contains definitions of:
   O_RDONLY (integer constant), open() */
#include <unistd.h>
int main(void) {
    int file_descriptor;
    file_descriptor = open("input.txt", O_RDONLY);
    if (file_descriptor < 0) {
        printf("error: %s\\n", strerror(errno));
        exit(1);
    }
    ...
    result = read(file_descriptor, ...);
    ...
}
```



# strace hello\_world (1)

strace — Linux tool to trace system calls  
run on assembly program we saw earlier:

```
$ strace -o trace.txt ./hello_world
$ cat trace.txt
execve("./hello_world", ["./hello_world"],
      0x7ffeedafd0a0 /* 28 vars */) = 0
write(1, "Hello, World!\n\0", 14)      = 14
exit(0)                                = ?
+++ exited with 0 +++
```



# strace hello\_world (2)

```
#include <stdio.h>
int main() { puts("Hello, World!"); }
```

---

when statically linked:

```
execve("./hello_world", [ "./hello_world" ], 0x7ffeb4127f70 /* 28 vars */)
    = 0
brk(NULL)
    = 0x22f8000
brk(0x22f91c0)
    = 0x22f91c0
arch_prctl(ARCH_SET_FS, 0x22f8880)
    = 0
uname({sysname="Linux", nodename="reiss-t3620", ...}) = 0
readlink("/proc/self/exe", "/u/cr4bd/spring2023/cs3130/slide"..., 4096)
    = 57
brk(0x231a1c0)
    = 0x231a1c0
brk(0x231b000)
    = 0x231b000
access("/etc/ld.so.nohwcap", F_OK)
    = -1 ENOENT (No such file or
                                directory)
fstat(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 4), ...}) = 0
write(1, "Hello, World!\n", 14)
    = 14
exit_group(0)
    = ?
+++ exited with 0 +++
```



# aside: what are those syscalls?

execve: run program

brk: allocate heap space

arch\_prctl(ARCH\_SET\_FS, ...): thread local storage pointer

may make more sense when we cover concurrency/parallelism later

uname: get system information

readlink of /proc/self/exe: get name of this program

access: can we access this file [in this case, a config file]?

fstat: get information about open file

exit\_group: variant of exit



# strace hello\_world (2)

```
#include <stdio.h>
int main() { puts("Hello, World!"); }
```

---

## when dynamically linked:

```
execve("./hello_world", [ "./hello_world" ], 0x7ffcfe91d540 /* 28 vars */)
                                     = 0
brk(NULL)                           = 0x55d6c351b000
...
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=196684, ...}) = 0
mmap(NULL, 196684, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f7a62dd3000
close(3)                             = 0
access("/etc/ld.so.nohwcap", F_OK)    = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0"... , 832) = 832
...
close(3)                             = 0
write(1, "Hello, World!\n", 14)       = 14
exit_group(0)                         = ?
+++ exited with 0 +++
```



# againframe(unixDesignHWNoLib)



# againframe(unixDesignHWLib)



# againframe(portalShouldnt)



# memory protection

modifying another program's memory?

## Program 1

```
long global = 42;
main() {
    printf("%p", &global);
    // 0x410010
    ...
    printf("%ld\n", global);
}
```

## Program 2

```
// while program 1 is in ...:
long *ptr;
ptr = (long*) 0x410010;
*ptr = 100;
```

What happens?

- A. 42 is printed
- B. 100 is printed
- C. program 1 crashes
- D. program 2 crashes
- E. something else




# program memory (two programs)




# program memory (two programs)

Program A




Program B




# program memory (two programs)



Program A



Program B



# program memory (two programs)



# address space

programs have *illusion of own memory*


called a program's *address space*



# address space

programs have *illusion of own memory*


called a program's *address space*



# address space

programs have *illusion of own memory*


called a program's *address space*



# address space

programs have *illusion of own memory*

called a program's *address space*



# againframe(progMem)



# againframe(addrSpace)



# address space mechanisms

topic after exceptions

called *virtual memory*

mapping called *page tables*

mapping part of what is changed in context switch



# program crashing?

what happens on processor when program crashes?

other program informed of crash to display message

use processor to run some other program

how does hardware do this?

would be complicated to tell about other programs, etc.

instead: hardware runs designated OS routine



# exceptions

recall: system calls — software asks OS for help

also cases where hardware asks OS for help

different triggers than system calls

but same mechanism as system calls:

- switch to kernel mode (if not already)

- call OS-designated function



# exceptions

recall: system calls — software asks OS for help

also cases where hardware asks OS for help

different triggers than system calls

but *same mechanism as system calls*:

- switch to kernel mode (if not already)

- call OS-designated function



# types of exceptions

system calls

- intentional — ask OS to do something

errors/events in programs

- memory not in address space (“Segmentation fault”)

- privileged instruction

- divide by zero, invalid instruction

- ...

external — I/O, etc.

- timer — configured by OS to run OS at certain time

- I/O devices — key presses, hard drives, networks, ...

- hardware is broken (e.g. memory parity error)



# types of exceptions

## *system calls*

intentional — ask OS to do something

## errors/events in programs

memory not in address space (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

## external — I/O, etc.

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

*errors/events in programs*

memory not in address space (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

external — I/O, etc.

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

memory not in address space (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

*external — I/O, etc.*

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

memory not in address space (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

external — I/O, etc.

*timer* — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

memory not in address space (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

external — I/O, etc.

timer — configured by OS to run OS at certain time

*I/O devices* — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

*memory not in address space* (“Segmentation fault”)

privileged instruction

divide by zero, invalid instruction

...

external — I/O, etc.

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

memory not in address space (“Segmentation fault”)

*privileged instruction*

divide by zero, invalid instruction

...

external — I/O, etc.

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# types of exceptions

system calls

intentional — ask OS to do something

errors/events in programs

memory not in address space (“Segmentation fault”)

privileged instruction

*divide by zero, invalid instruction*

...

external — I/O, etc.

timer — configured by OS to run OS at certain time

I/O devices — key presses, hard drives, networks, ...

hardware is broken (e.g. memory parity error)



# againframe(exceptTypesN)




# againframe(portalShouldnt)




# againframe(exceptTypesN)



# exceptions [Venn diagram]





# general exception process



# time multiplexing




# time multiplexing



= operating system

# time multiplexing



= operating system

exception happens

return from exception

# switching programs

OS starts running somehow

some sort of exception

saves old registers + program counter + address mapping

(optimization: could omit when program crashing/exiting)


sets new registers + address mapping, jumps to new  
program counter

called *context switch*


saved information called *context*



# contexts (A running)



# contexts (B running)



# threads

thread = illusion of own processor

own register values

own program counter value

actual implementation:

many threads sharing one processor

problem: where are register/program counter values  
when thread not active on processor?



# againframe(exceptTypesN)



# exception patterns with I/O (1)

input — available now:

exception: device says “I have input now”

handler: OS stores input for later

exception (syscall): program says “I want to read input”

handler: OS returns that input

input — not available now:

exception (syscall): program says “I want to read input”

handler: OS runs other things (context switch)

exception: device says “I have input now”

handler: OS retrieves input

handler: (possibly) OS switches back to program that wanted it



# exception patterns with I/O (2)

output — ready now:

exception (syscall): program says “I want to output this”

handler: OS sends output to device

output — not ready now

exception (syscall): program says “I want to output”


handler: OS realizes device can’t accept output yet  
(other things happen)

exception: device says “I’m ready for output now”


handler: OS sends output requested earlier




# keyboard input timeline



# context switches and exceptions



 = operating system

program execution interleaves with OS

OS might context switch when it is run, might not

# emptyloop lab (1)

```
long last_time = GetTime();  
while (NotDone()) {  
    long current_time = GetTime();  
    RecordDelta(current_time - last_time);  
    last_time = current_time;  
}
```

naive computer model: RecordDelta called with roughly same number each time



# emptyloop lab (2)

```
long last_time = GetTime();
while (NotDone()) {
    long current_time = GetTime();
    /* maybe OS handles I/O here? */
    /* maybe OS runs ssh for a bit here? */
    /* ... */
    RecordDelta(current_time - last_time);
    last_time = current_time;
}
```

will see spikes in recorded times from exceptions  
can infer what system is doing



# emptyloop lab (3)

upcoming lab: we'll supply program that times in loop

also will look at Linux counters for asynchronous exceptions

should be able to observe when OS/other programs run

(and about how long)



# review: definitions

exception: hardware calls OS specified routine

many possible reasons

system calls: type of exception

context switch: OS switches to another thread

by saving old register values + loading new ones

part of OS routine run by exception



# which of these require exceptions? context switches?

- A. program calls a function in the standard library
- B. program writes a file to disk
- C. program A goes to sleep, letting program B run
- D. program exits
- E. program returns from one function to another function
- F. program pops a value from the stack



# which require exceptions [answers] (1)

A. program calls a function in the standard library

no (same as other functions in program); many standard library functions make no system calls (and do not otherwise trigger exceptions — for example `strlen`, `pow`; also if we consider the calling of a function just the `call` instruction, then the library functions that do make system calls won't do so until later)

B. program writes a file to disk

yes (requires kernel mode only operations)

C.  program A goes to sleep, letting program B run

# which require exceptions [answer] (2)

D. program exits

yes (requires switching to another program, which requires accessing OS data + other program's memory)

E. program returns from one function to another function

no

F. program pops a value from the stack

no



# which require context switches [answer]

no: A. program calls a function in the standard library

no: B. program writes a file to disk

(but might be done if program needs to wait for disk and other things could be run while it does)

yes: C. program A goes to sleep, letting program B run

yes: D. program exits

no: E. program returns from one function to another function

no: F. program pops a value from the stack



# terms for exceptions

terms for exceptions aren't standardized

our readings use one set of terms

interrupts = externally-triggered

faults = error/event in program

trap = intentionally triggered

all these terms appear differently elsewhere



# The Process

*process* = thread(s) + address space

illusion of *dedicated machine*:

thread = illusion of own CPU

(process could have multiple threads – with independent registers)

address space = illusion of own memory

