

CSO2 (CS3130)

# themes

automating building software

- libraries, taking advantage of incremental compilation

sharing machines

- multiple users/programs on one system

parallelism and concurrency

- doing two+ things at once

under the hood of sockets

- layered design of networks

- implementing secure communication

under the hood of fast processors

- caching, (hidden) parallelism, avoiding idle time

# themes

## automating building software

- libraries, taking advantage of incremental compilation

## sharing machines

- multiple users/programs on one system

## parallelism and concurrency

- doing two+ things at once

## under the hood of sockets

- layered design of networks
- implementing secure communication

## under the hood of fast processors

- caching, (hidden) parallelism, avoiding idle time

# make

```
$ ./foo.exe
```

```
...
```

```
...
```

```
$ edit readline.c
```

```
$ make
```

```
clang -g -O -Wall -c readline.c -o readline.o
```

```
ar rcs terminal.o readline.o libreadline.a
```

```
clang -o foo.exe foo.o foo-utility.o -L. -lreadline
```

```
$
```

# themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design of networks

implementing secure communication

under the hood of fast processors

caching, (hidden) parallelism, avoiding idle time

# address translation



# address translation



# address translation



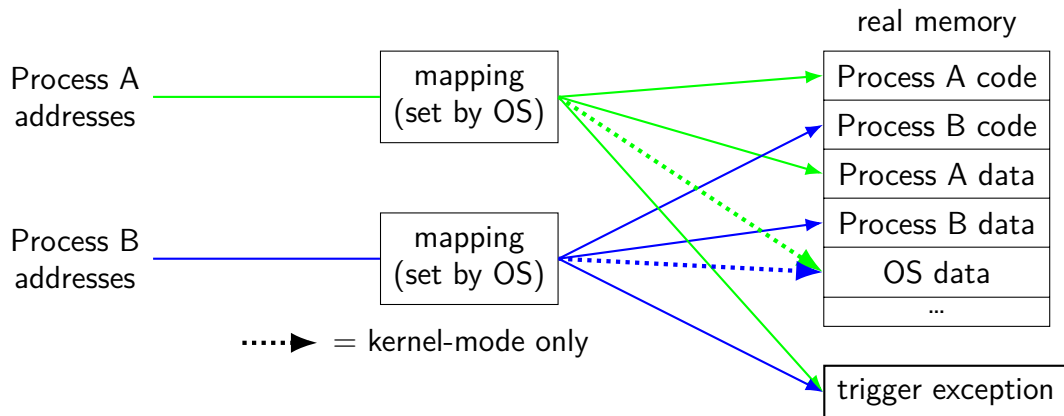


# address translation



# address spaces

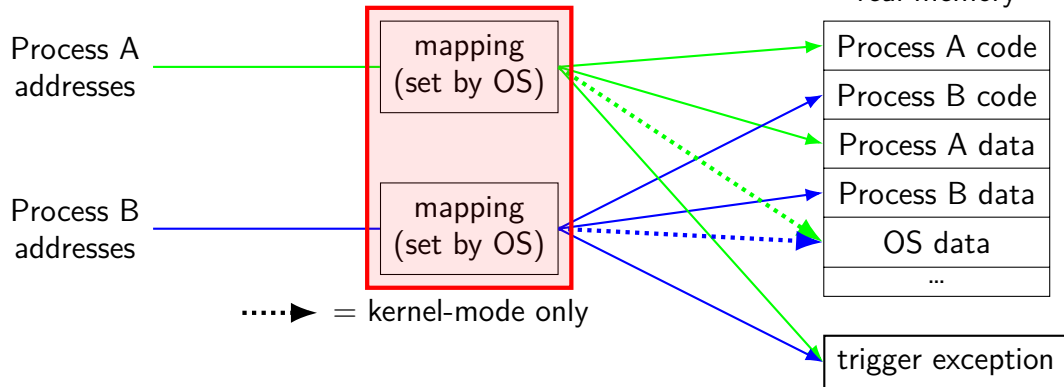
illusion of **dedicated memory**



# address spaces

illusion of **dedicated memory**

chose one during context switch



# themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

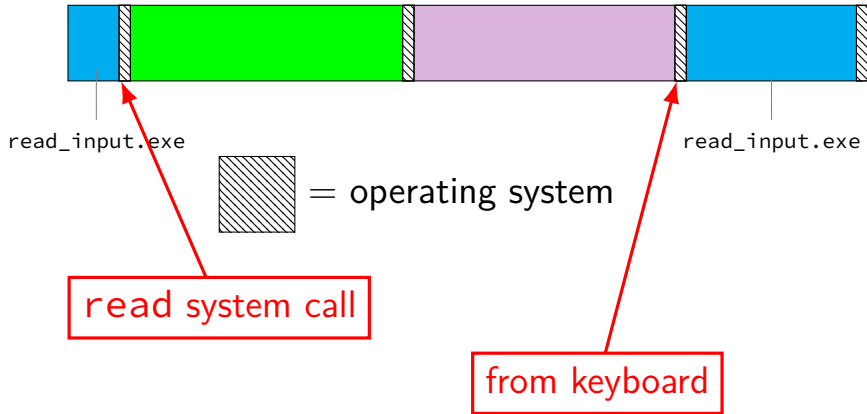
layered design of networks

implementing secure communication

under the hood of fast processors

caching, (hidden) parallelism, avoiding idle time

# keyboard input timeline



# time multiplexing

processor:



# time multiplexing



...

```
call get_time
```

```
// whatever get_time does
```

```
movq %rax, %rbp
```

———— million cycle delay ————

```
call get_time
```

```
// whatever get_time does
```

```
subq %rbp, %rax
```

...

# time multiplexing



...

```
call get_time
```

```
// whatever get_time does
```

```
movq %rax, %rbp
```

———— million cycle delay ————

```
call get_time
```

```
// whatever get_time does
```

```
subq %rbp, %rax
```

...



# multiple cores+threads

core 1:



core 2:




multiple cores? each core still divided up

# multiple cores+threads

core 1:



core 2:



one program with multiple *threads*

# themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design of networks

implementing secure communication

under the hood of fast processors

caching, (hidden) parallelism, avoiding idle time

# permissions

```
$ ls /u/other/secret
```

```
ls: cannot open directory '/u/other/secret': Permission denied
```

```
$ shutdown
```

```
shutdown: Permission denied
```

# themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design of networks

implementing secure communication

under the hood of fast processors

caching, (hidden) parallelism, avoiding idle time

# layers

application	HTTP, SSH, SMTP, ...	application-defined meanings
transport	TCP, UDP, ...	reach      correct      program, reliability/streams
network	IPv4, IPv6, ...	reach      correct      machine (across networks)
link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
physical	...	encode bits for wire/radio

# names and addresses

name	address
logical identifier	location/how to locate
variable counter	memory address 0x7FFF9430
DNS name www.virginia.edu	IPv4 address 128.143.22.36
DNS name mail.google.com	IPv4 address 216.58.217.69
DNS name mail.google.com	IPv6 address 2607:f8b0:4004:80b::2005
DNS name reiss-t3620.cs.virginia.edu	IPv4 address 128.143.67.91
DNS name reiss-t3620.cs.virginia.edu	MAC address 18:66:da:2e:7f:da
service name https	port number 443
service name ssh	port number 22

# secure communication?

how do you know who your socket is to?

who can read what's on the socket?

what can you do to restrict this?



# themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design of networks

implementing secure communication

under the hood of fast processors

caching, (hidden) parallelism, avoiding idle time

# 2004 CPU

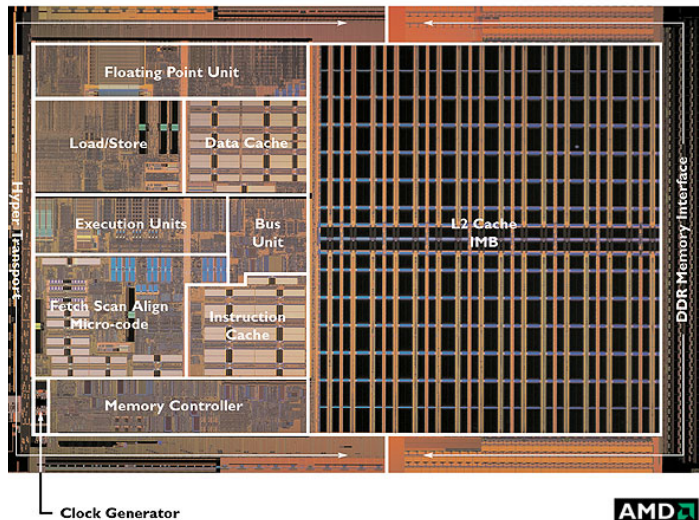


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)

# 2004 CPU

▲ Registers

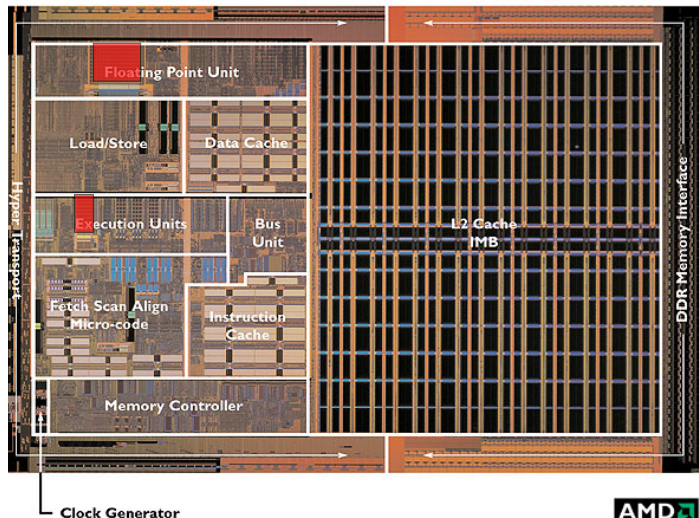


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)

# 2004 CPU

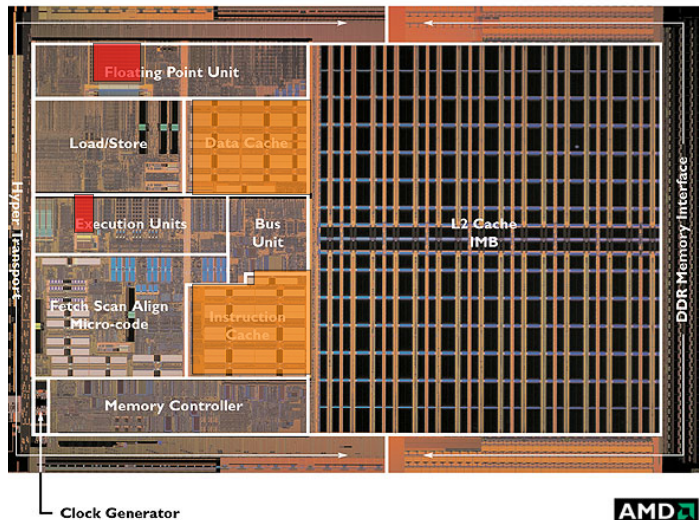


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)

# 2004 CPU

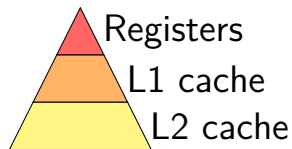
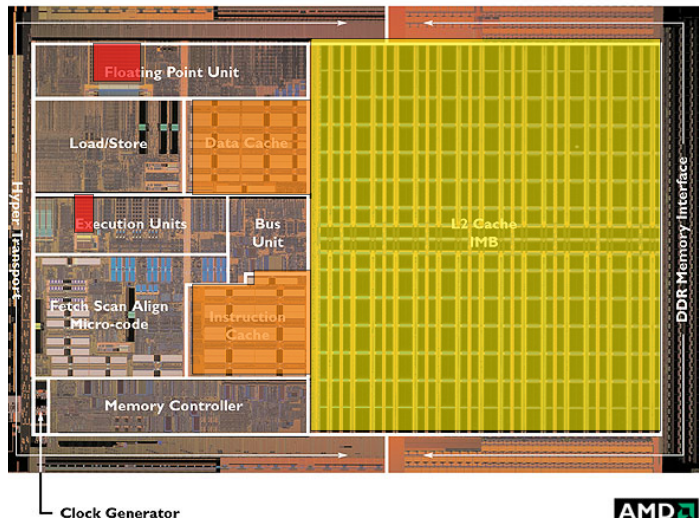


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)

# 2004 CPU



Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)

# 2004 CPU

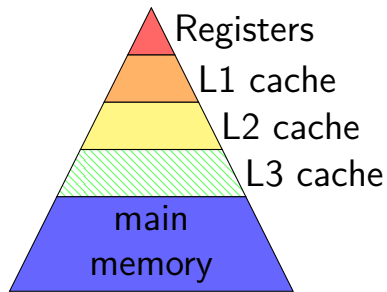


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)



# 2004 CPU


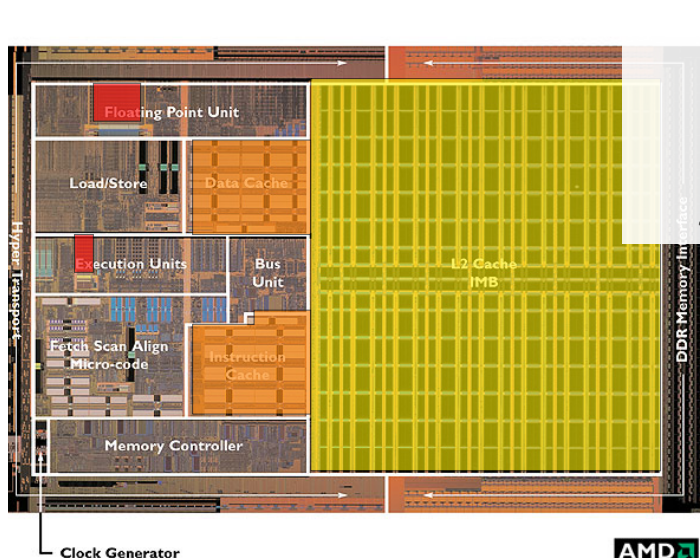


Image: approx 2004 AMD press image of Opteron die;  
approx register location via [chip-architect.org](http://chip-architect.org) (Hans de Vries)



## some performance examples

example1:

```
    movq $1000000000000, %rax
loop1:
    addq %rbx, %rcx
    decq %rax
    jge loop1
    ret
```

about 30B instructions

my desktop: approx 2.65 sec

example2:

```
    movq $1000000000000, %rax
loop2:
    addq %rbx, %rcx
    addq %r8, %r9
    decq %rax
    jge loop2
    ret
```

about 40B instructions

my desktop: approx 2.65 sec

# some performance examples

example1:

```
    movq $1000000000000, %rax
loop1:
    addq %rbx, %rcx
    decq %rax
    jge loop1
    ret
```

about 30B instructions

my desktop: approx 2.65 sec

example2:

```
    movq $1000000000000, %rax
loop2:
    addq %rbx, %rcx
    addq %r8, %r9
    decq %rax
    jge loop2
    ret
```

about 40B instructions

my desktop: approx 2.65 sec

## C exercise

```
int array[4] = {10,20,30,40};  
int *p;  
p = &array[0];  
p += 2;  
p[1] += 1;
```

array =

- |                             |                  |
|-----------------------------|------------------|
| A. compile or runtime error | B. {10,20,30,41} |
| C. {10,20,32,41}            | D. {10,21,30,40} |
| E. {12,21,30,40}            | F. none of these |

# some avenues for review

review CSO1 stuff

labs 9–12 (of last Spring)

<https://www.cs.virginia.edu/~jh2jf/courses/cs2130/spring2023/>

exercises we've used in the past:

implement strsep library function

implement conversion from dynamic array to linked list

## some pointer stuff

0x040

0x038

0x030

0x028

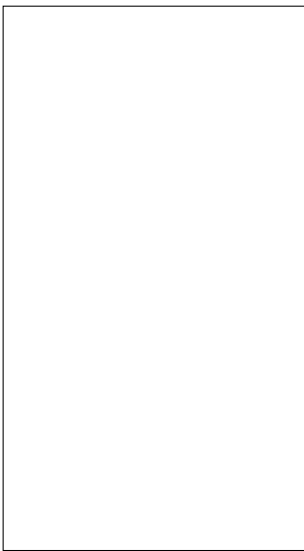
0x020

0x018

0x010

0x008

0x000



```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

## some pointer stuff

0x040

0x038

0x030

0x028

0x020

0x018

0x010

0x008

0x000

array[2]: 0x67
array[1]: 0x45
array[0]: 0x12
single: 0x78
ptr = ???

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

## some pointer stuff

0x040	
0x038	
0x030	array[2]: 0x67
	array[1]: 0x45
	array[0]: 0x12
0x028	single: 0x78
0x020	ptr = ???
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

~~\*ptr = 0xAB;~~ compile error

## some pointer stuff

0x040	
0x038	
0x030	array[2]: 0x67
	array[1]: 0x45
	array[0]: 0x12
0x028	single: 0x78
	ptr: 0x28
0x020	
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

```
ptr = &single;
```

```
ptr = (int*) 0x28;    addr. of single
```



## some pointer stuff

0x040	
0x038	
0x030	array[2]: 0x67
	array[1]: 0x45
	array[0]: 0x12
0x028	single: 0x78
	ptr: 0x28
0x020	
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

```
ptr = &single;  
ptr = (int*) 0x28;    addr. of single
```

~~ptr = 0x28; compile error~~

~~ptr = (int\*) single;~~  
pointer to unknown place

## some pointer stuff

0x040	
0x038	array[2]: 0x67
0x030	array[1]: 0x45
	array[0]: 0x12
0x028	single: 0xFF
	ptr: 0x28
0x020	
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;  
ptr = &single;
```

```
*ptr = 0xFF;
```

## some pointer stuff

0x040	
0x038	
0x030	array[2]: 0x67
	array[1]: 0x45
	array[0]: 0x12
0x028	single: 0x78
	ptr: 0x2C
0x020	
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

```
ptr = array;  
ptr = &array[0];  
ptr = (int*) 0x2C;
```

## some pointer stuff

0x040

0x038

0x030

0x028

0x020

0x018

0x010

0x008

0x000

array[2]: 0x67
array[1]: 0x45
array[0]: 0x12
single: 0x78
ptr: 0x2C

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

```
ptr = array;  
ptr = &array[0];  
ptr = (int*) 0x2C;
```

~~ptr = array[0]; compile error~~

~~ptr = (int\*) array[0];~~

pointer to unknown place

## some pointer stuff

0x040

0x038

0x030

0x028

0x020

0x018

0x010

0x008

0x000

array[2]: 0xFF
array[1]: 0x45
array[0]: 0x12
single: 0x78
ptr: 0x2C

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;  
ptr = &array[0];
```

```
ptr[2] = 0xFF;  
*(ptr + 2) = 0xFF;
```

```
int *temp1; temp1 = ptr + 2;  
*temp1 = 0xFF;
```

```
int *temp2; temp2 = &ptr[2];  
*temp2 = 0xFF;
```

## some pointer stuff

0x040	
0x038	
0x030	array[2]: 0x67
	array[1]: 0x45
	array[0]: 0x12
0x028	single: ...
	ptr: 0x2C
0x020	
0x018	
0x010	
0x008	
0x000	

```
int array[3]={0x12,0x45,0x67};  
int single = 0x78;  
int *ptr;
```

```
void change_arg(int *x) {  
    *x = compute_some_value();  
}  
...  
change_arg(&single);
```

# backup slides