

CS 3130 intro

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

make

```
$ ./foo.exe
```

```
...
```

```
...
```

```
$ edit readline.c
```

```
$ make
```

```
clang -g -O -Wall -c readline.c -o readline.o
```

```
ar rcs terminal.o readline.o libreadline.a
```

```
clang -o foo.exe foo.o foo-utility.o -L. -lreadline
```

```
$
```

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

address translation



address translation



address translation

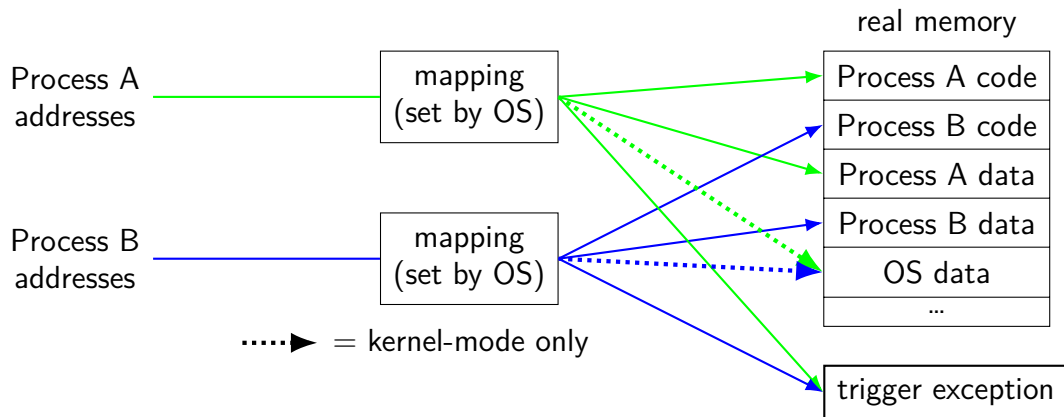


address translation



address spaces

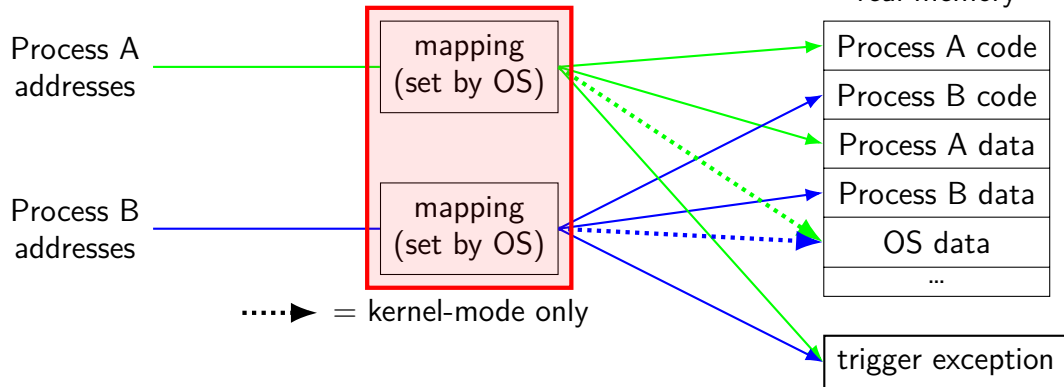
illusion of **dedicated memory**



address spaces

illusion of **dedicated memory**

chose one during context switch



themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

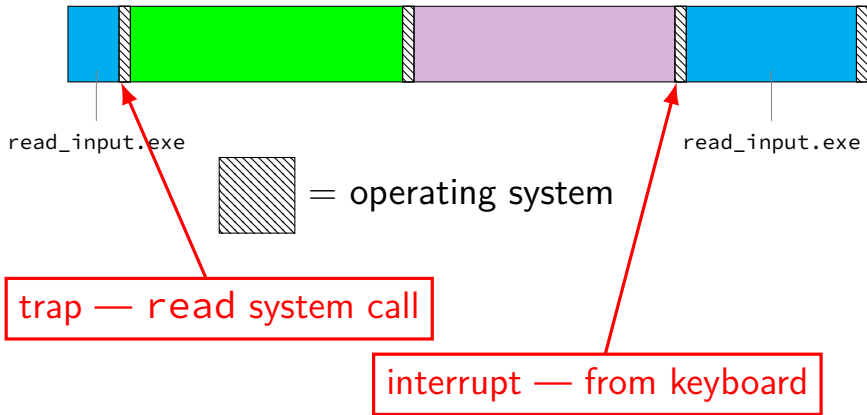
under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

keyboard input timeline



time multiplexing

processor:



time multiplexing

processor:



...

```
call get_time
```

```
// whatever get_time does
```

```
movq %rax, %rbp
```

———— million cycle delay ————

```
call get_time
```

```
// whatever get_time does
```

```
subq %rbp, %rax
```

...

time multiplexing



...

```
call get_time
```

```
// whatever get_time does
```

```
movq %rax, %rbp
```

———— million cycle delay ————

```
call get_time
```

```
// whatever get_time does
```

```
subq %rbp, %rax
```

...

multiple cores+threads

core 1:



core 2:




multiple cores? each core still divided up

multiple cores+threads

core 1:



core 2:



one program with multiple *threads*

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

permissions

```
$ ls /u/other/secret
```

```
ls: cannot open directory '/u/other/secret': Permission denied
```

```
$ shutdown
```

```
shutdown: Permission denied
```

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

layers

| | | |
|-------------|----------------------|---|
| application | HTTP, SSH, SMTP, ... | application-defined meanings |
| transport | TCP, UDP, ... | reach correct program, reliability/streams |
| network | IPv4, IPv6, ... | reach correct machine (across networks) |
| link | Ethernet, Wi-Fi, ... | coordinate shared wire/radio |
| physical | ... | encode bits for wire/radio |

more than four layers?

sometimes more layers above 'application'

e.g. HTTPS:

HTTP (app layer) on TLS (another app layer) on TCP (network) on ...

e.g. DNS over HTTPS:

DNS (app layer) on HTTP on on TLS on TCP on ...

e.g. SFTP:

SFTP (app layer??) on SSH (another app layer) on TCP on ...

e.g. HTTP over OpenVPN:

HTTP on TCP on IP on OpenVPN on UDP on different IP on ...

names and addresses

| name | address |
|--------------------------------------|---------------------------------|
| logical identifier | location/how to locate |
| variable counter | memory address 0x7FFF9430 |
| DNS name www.virginia.edu | IPv4 address 128.143.22.36 |
| DNS name mail.google.com | IPv4 address 216.58.217.69 |
| DNS name mail.google.com | IPv6 address 2607:f8b0:4004:80b |
| DNS name reiss-t3620.cs.virginia.edu | IPv4 address 128.143.67.91 |
| DNS name reiss-t3620.cs.virginia.edu | MAC address 18:66:da:2e:7f |
| service name https | port number 443 |
| service name ssh | port number 22 |

secure communication?

how do you know who your socket is to?

who can read what's on the socket?

what can you do to restrict this?

themes

automating building software

libraries, taking advantage of incremental compilation

sharing machines

multiple users/programs on one system

parallelism and concurrency

doing two+ things at once

under the hood of sockets

layered design on networks

under the hood of fast processors

caching and (hidden) parallelism

2004 CPU

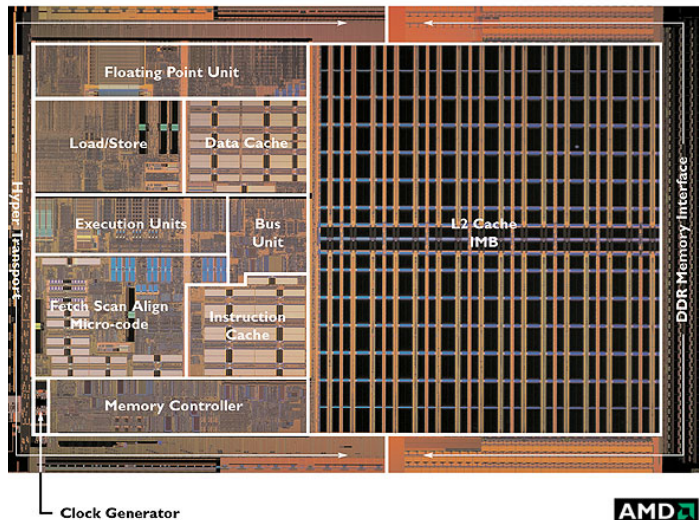


Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU

▲ Registers

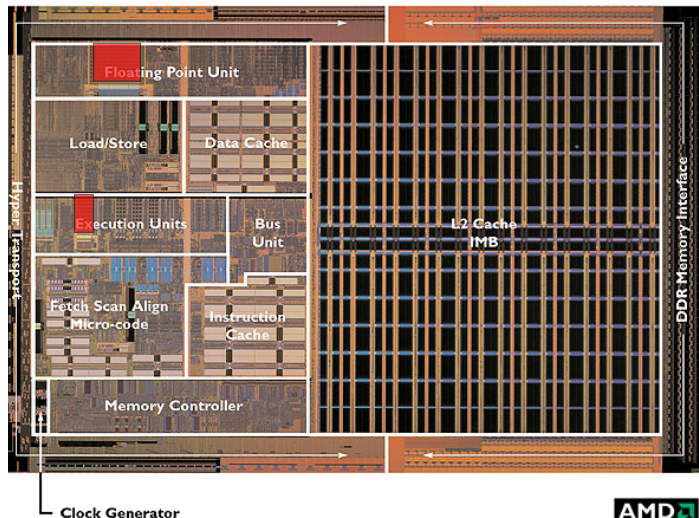


Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU

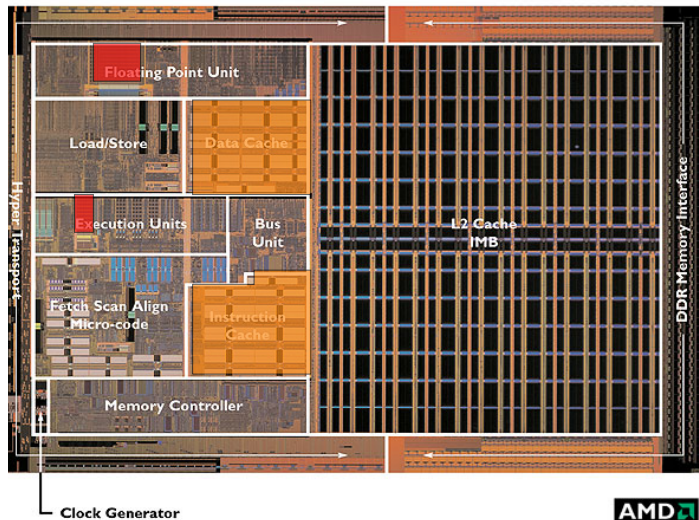


Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU



Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU

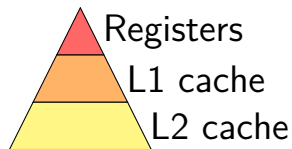


Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU

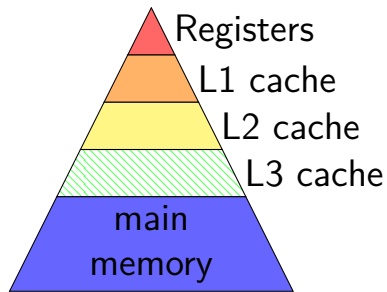
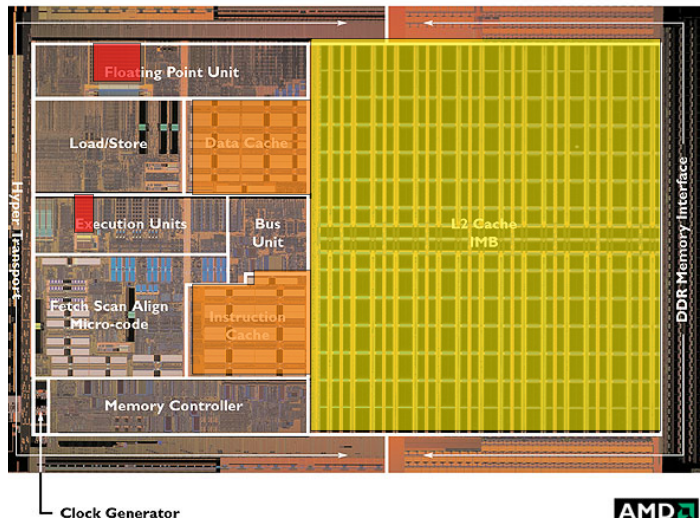
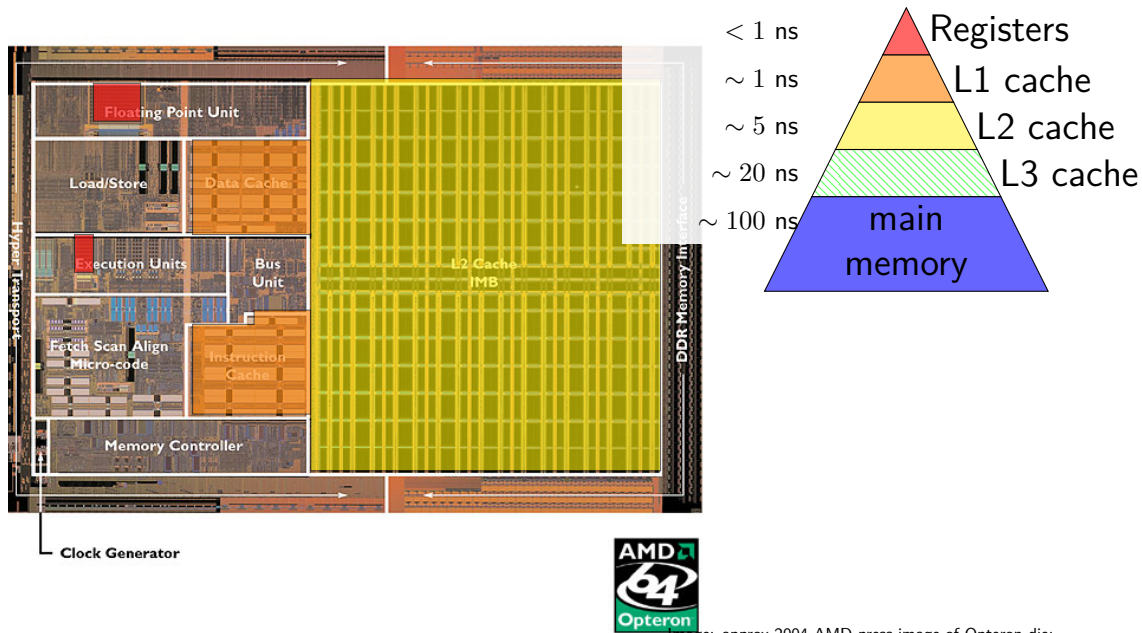


Image: approx 2004 AMD press image of Opteron die;
approx register location via chip-architect.org (Hans de Vries)

2004 CPU



themes

automating building software

- libraries, taking advantage of incremental compilation

sharing machines

- multiple users/programs on one system

parallelism and concurrency

- doing two+ things at once

under the hood of sockets

- layered design on networks

under the hood of fast processors

- caching and (hidden) parallelism

some performance examples

example1:

```
    movq $1000000000000, %rax
loop1:
    addq %rbx, %rcx
    decq %rax
    jge loop1
    ret
```

about 30B instructions

my desktop: approx 2.65 sec

example2:

```
    movq $1000000000000, %rax
loop2:
    addq %rbx, %rcx
    addq %r8, %r9
    decq %rax
    jge loop2
    ret
```

about 40B instructions

my desktop: approx 2.65 sec

some performance examples

example1:

```
    movq $1000000000000, %rax
loop1:
    addq %rbx, %rcx
    decq %rax
    jge loop1
    ret
```

about 30B instructions

my desktop: approx 2.65 sec

example2:

```
    movq $1000000000000, %rax
loop2:
    addq %rbx, %rcx
    addq %r8, %r9
    decq %rax
    jge loop2
    ret
```

about 40B instructions

my desktop: approx 2.65 sec

logistics

labs

attend lab in person and get checked off by TA, *or*

(most labs) submit something to submission site and we'll grade it

submit to submission site? don't care if you attend the lab

more strict about submissions without checkoffs being complete/correct

(can't tell how much time you actually spent)

in-person lab checkoff of incomplete lab at least 50% credit

some labs will basically require attendance

or contact me for other arrangements if you can't (sick, etc.)

logistically won't work otherwise — e.g. code review

lab collaboration and submissions

please collaborate on labs!

when working with others on lab and submitting code files

please indicate who you worked with in those files
via comment or similar

quizzes

released evening after Thursday lecture
starting *next* week

due 15 minutes before lecture on Tuesdays

about lecture and/or lab from the prior week

4–6 questions

individual, open book, open notes, open Internet

okay: looking up resources/tutorials/etc.

not okay: asking Stack Overflow the quiz question

not okay: IMing your friend the quiz question

asking about quiz questions

I and the TAs won't answer quiz questions...

but we will answer questions about the lecture material, etc.

(and TAs (not you) are responsible for knowing
what they can't answer

but we'd prefer you don't try to test those limits)

homeworks

several homework assignments

done individually

due before a week's first lab

exams

1 final exam

no midterms — instead:

- quizzes count a lot

- slightly more homework/lab than pilot

development enviroment

official: department machines via SSH or NX (remote desktop)

you can also use your own machines, but...

we will test your code on x86-64 Linux

I haven't checked assignments on a Windows or OS X machine

getting help

office hours — calendar will be posted on website

mix of in-person and remote, indicated on calendar

remote OH will use Discord + online queue

in-person OH may or may not — indicated on whiteboard, probably

Piazza

use private questions if homework code, etc.

emailing me (preferably with '3130' in subject)

late policy

no late quizzes

one quiz dropped (unconditionally)

90% credit for 0–48 hours late homeworks

80% credit for 48–72 hours late homeworks

for labs that allow submission only, same policy as homeworks

lab submission due time is 11:59pm

for other labs, policy on a lab-by-lab basis

excused lateness

special circumstances?

illness, emergency, etc.

contact me, we'll figure something out

please don't attend lab/etc. sick!


attendance

I won't take attendance in lecture

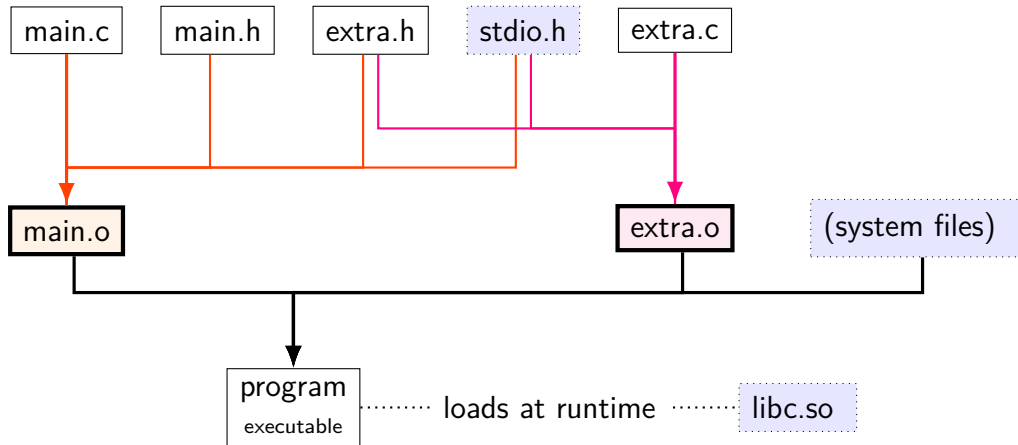
I will attempt to have lecture recordings

sometimes there may be issues with the recording

files in building C programs [dynamic linking]




files in building C programs [dynamic linking]




```
clang -c main.c  
clang -c extra.c
```

files in building C programs [dynamic linking]



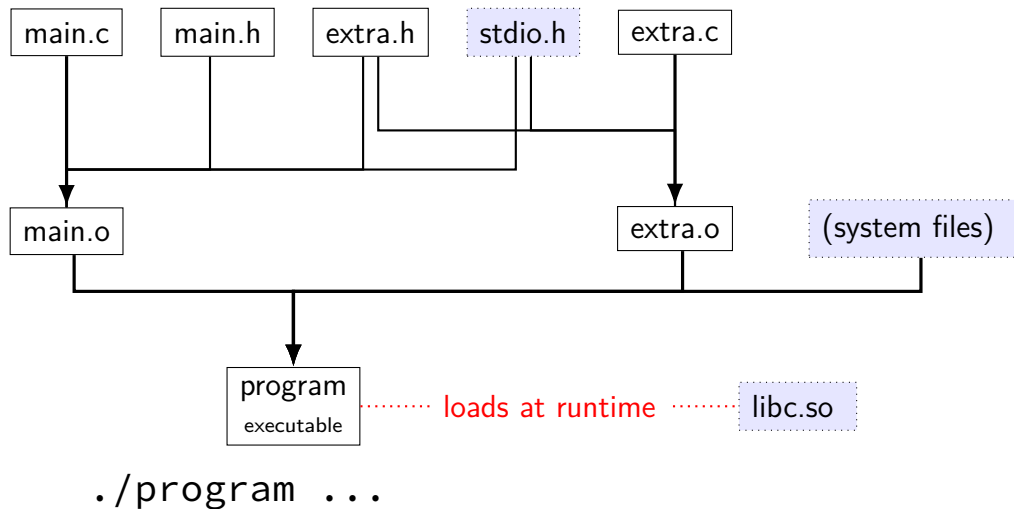
```
clang -S -c main.c  
clang -S -c extra.c
```

files in building C programs [dynamic linking]

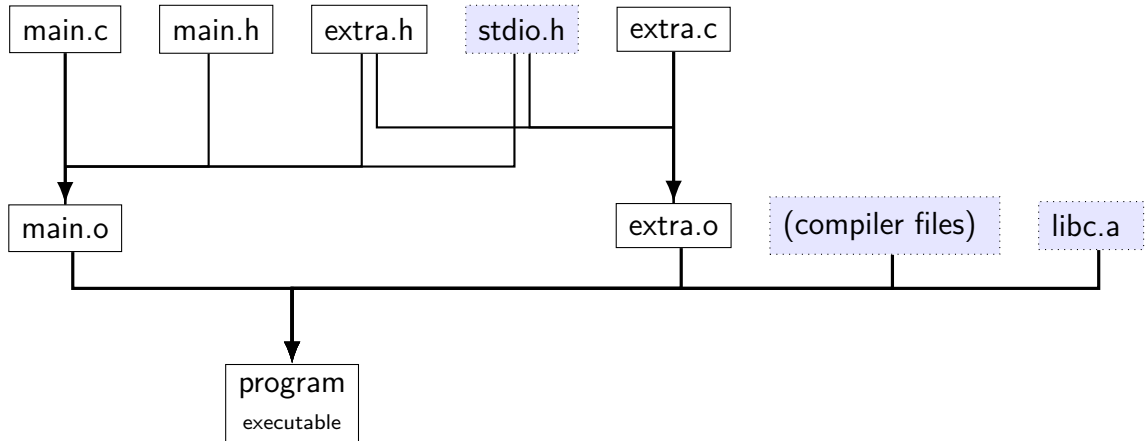


```
clang -o program main.o extra.o
```

files in building C programs [dynamic linking]



files in building C programs [static linking]



file extensions

| name | | |
|--------|-----------|--|
| .c | | C source code |
| .h | | C header file |
| .s | (or .asm) | assembly file |
| .o | (or .obj) | object file (binary of assembly) |
| (none) | (or .exe) | executable file |
| .a | (or .lib) | statically linked library [collection of .o files] |
| .so | (or .dll) | dynamically linked library ['shared object'] |

static libraries

Unix-like *static* libraries: libfoo.a

internally: archive of .o files with index

create: `ar rcs file1.o file2.o ...`

use: `cc ... -o program -L/path/to/lib ...-lfoo`

`cc` could be `clang`, `gcc`, `clang++`, `g++`, etc.

`-L/path/to/lib` not needed if in standard location

shared libraries

Linux *shared* libraries: libfoo.so

create:

compile .o files with -fPIC (position independent code)

then: `cc -shared ... -o libfoo.so`

use: `cc ...-o program -L/path/to/lib ...-lfoo`

finding shared libraries

```
cc ...-o program -L/path/to/lib ...-lfoo
```

on Linux: `/path/to/lib` only used to create program
program contains `libfoo.so` *without full path*

Linux default: `libfoo.so` expected to be in `/usr/lib`, `/lib`, and other 'standard' locations

possible overrides:

`LD_LIBRARY_PATH` environment variable
paths specified with `-Wl,-rpath=/path/to/lib` when creating executable

exercise (incremental compilation)

program built from main.c + extra.c

main.c, extra.c both include extra.h, stdio.h

```
clang -c main.c           # command 1
clang -c extra.c          # command 2
clang -o program main.o extra.o # command 3
```

What commands need to be rerun if...

Question A: ...main.c changes?

Question B: ...extra.h changes?

backup slides