



# changelog

17 Feb 2023: memory access example: shift things so address labels are not off from numbers shown

# last time (1)

assignment Q&A

multi-level page table lookup

unreliable networks

“best effort” model of the internet

- limited message size

- sometimes messages lost

- sometimes messages delayed/reordered

- sometimes messages corrupted

## last time (2)

sequence numbers

acknowledgments

- someone needs to resend after timeout
- can get lost, that's okay

checksums

# some themes in anonymous feedback

pageable difficulty

lab difficulty

quizzes: how many/etc.

## Quiz Q3

- 1 first-level page table with
    - a valid entry pointing to a second-level page table with 512 valid entries
    - a valid entry pointing to a second-level page table with (1000-512) valid entries and a few invalid entries and 510 invalid entries
- three 4096-byte page tables

## Quiz Q4

$$\begin{aligned} 0x120008 &= \text{PTBR} + \text{VPN part 1} \times \text{PTE size} = \\ 0x1200000 + \text{VPN part 1} \times 8 &\rightarrow \text{VPN part 1} = 1 \end{aligned}$$

$$\begin{aligned} 0x123040 &= \\ \text{PPN from 1st level} \times \text{page size} + \text{VPN part 2} \times \text{PTE size} &\rightarrow \\ \text{VPN part 2} &= 8 \end{aligned}$$

$$\begin{aligned} 0x6010 &= \\ \text{PPN from 2nd level} \times \text{page size} + \text{page offset} &\rightarrow \text{page offset} = 0x10 \end{aligned}$$

## Quiz Q5

“It then runs a function, whose machine code is loaded at addresses 0x2040-0x2072, which writes 3 8-byte values to the stack at addresses 0xFFFF8, 0xFFFF0, and 0xFFE8.”

page at 0x2000–0x2FFF

- code loaded on first instruction's page fault

- can't tell processor about only part of page being loaded

page at 0xF000–0xFFFF

- whole page of stack allocated on first access



# HW difficulty

“...I feel like several components of the assignment we have not fully learned and some we just learned about in lecture today. Additionally, I think while a checkpoint is a reasonable idea, we could all benefit from the extra time and just have the first two parts be due next week. I have been in office hours the last two days and it seems like barely any students know what is going on.”

“While the quiz made sense and was related to the lectures and readings, this homework assignment has a lot of things that you need to rely on TA's or word of mouth for. For example, how would we know that we need to memset after `posix_memalign` if we don't even know how to look that up...”

“I feel like the content of the lectures is too far removed from what we are asked to do in the homeworks....”

# mistakes I made with homework (1)

- overestimated C familiarity from CSO1

  - a lot of problems from C pointer issues

  - fails in ways that are not intuitive, especially if you aren't checking every step

  - why I assumed understanding manpage for `posix_memalign` was not big deal

  - future: warmup assignment should probably review C pointer stuff somehow

  - b/c of this, put halfway point of assignment at wrong place

- in future semesters, need to plan more lecture time for virtual memory

## mistakes I made with homework (2)

some things in writeup are/were too easy to miss

- page table entry format

- physical page number v physical address

- what things need to be allocated

need more structure re: testing

- students just using code in assignment + autograder was not the intention

- seems like (based on submissions) many students writing a lot of code before testing it, rather than testing in small pieces

# pagetable grading

submission 1 (25% of normal homework)

- 32% LEVELS = 1 reasonable attempt

- 64% reasonable attempt on two other items

- 3% code style

submission 2 (25% of normal homework)

- 50% everything present

- 40% LEVELS = 1 functionality

- 10% LEVELS > 1 functionality

submission 3 (200% of normal homework)

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

size\_t x = 0x50;  
~~\*\*x~~ (compile-time error)

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

~~\*x~~ (compile-time error)

```
size_t *ptr;
```

```
ptr = (size_t *) x;
```

```
*ptr == 0xABCDEF
```

```
*((size_t *) x) == 0xABCDEF
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

~~x[2]~~ (compile-time error)

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

```
x[2] (compile-time error)
```

```
size_t addr = x + 16;  
size_t *ptr;  
ptr = (size_t *) addr;  
*ptr == 0x456789
```

```
size_t *ptr;  
ptr = (size_t *) x;  
ptr[2] == 0x456789
```



## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;  
void change_arg(size_t *arg) {  
    *arg = 0xFFFF;  
}
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0xFFFF
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

```
void change_arg(size_t *arg) {  
    *arg = 0xFFFF;  
}
```

```
change_arg(&x);
```

```
change_arg((size_t*) 0x30);
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	? = 0xFFFF
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

```
void change_arg(size_t *arg) {  
    *arg = 0xFFFF;  
}
```

```
change_arg(&x + 1);  
change_arg((size_t*) 0x38);
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xFFFF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
size_t x = 0x50;
```

```
void change_arg(size_t *arg) {  
    *arg = 0xFFFF;  
}
```

```
change_arg((size_t *) x);  
change_arg((size_t *) 0x50);
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	
0x030	x = 0xFFFF
0x020	
0x010	
0x000	

```
void *x = (void *) 0x50
```

```
void change_arg(void **arg) {  
    *arg = (void *) 0xFFFF;  
}
```

```
change_arg((void **) &x);  
change_arg((void **) 0x30);
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xABCDEF
0x040	? = 0xFFFF
0x030	x = 0x50
0x020	
0x010	
0x000	

```
void *x = (void *) 0x50
```

```
void change_arg(void **arg) {  
    *arg = (void *) 0xFFFF;  
}
```

```
change_arg(&x + 1);  
change_arg((void **) 0x38);
```

## some pointer stuff

0x080	
0x070	
0x060	0x456789
	0x123456
0x050	0xFFFF
0x040	
0x030	x = 0x50
0x020	
0x010	
0x000	

```
void *x = (void *) 0x50
```

```
void change_arg(void **arg) {  
    *arg = (void *) 0xFFFF;  
}
```

```
change_arg((void **) x);  
change_arg((void **) 0x50);
```

# address/page table entry format

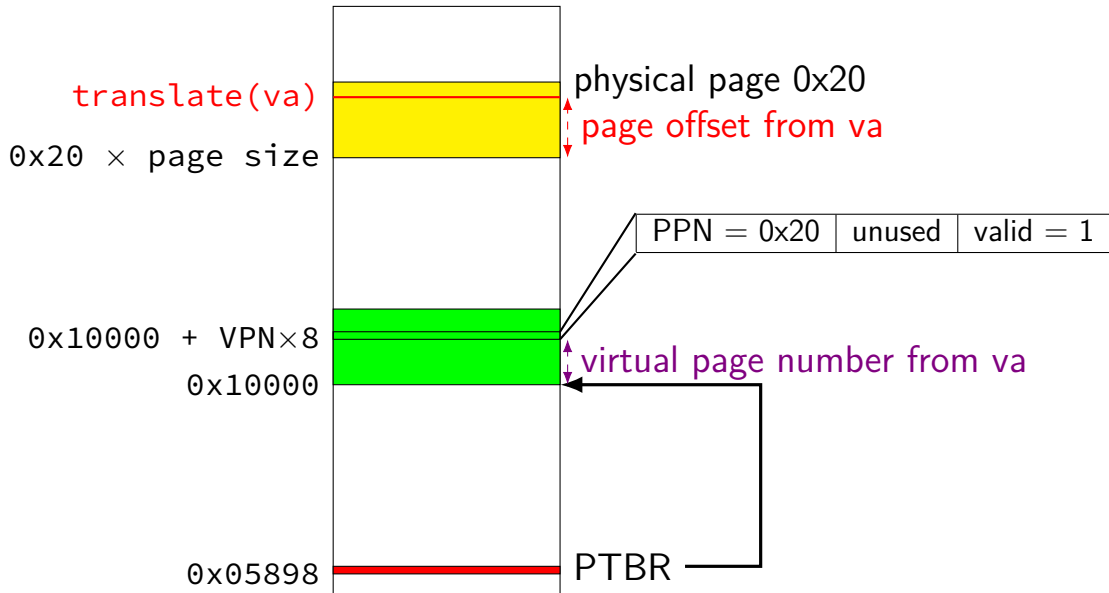
(with POBITS=12, LEVELS=1)

	bits 63-21	bits 20-12	bits 11-1	bit 0
page table entry	physical page number		unused	valid bit
virtual address	unused	virtual page number	page offset	
physical address	physical page number		page offset	


in assignment: value from `posix_memalign` = physical address



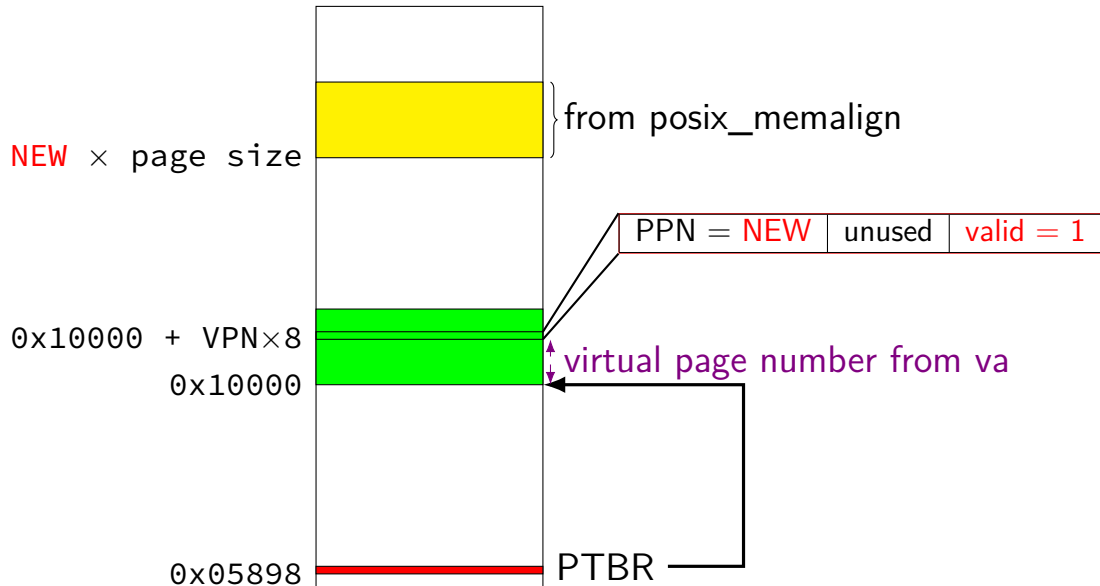
# pa = translate(va)



# page\_allocate(va)



# page\_allocate(va)



## next week's lab

code review your submission 2 with other students

must be in person!

can't attend lab? talk to me!

use the feedback to improve your submission 3

## lab difficulty

“I wish we could at least get more explanation for what is going on in the networking lab. I understood Tuesday’s lecture enough to at least get the concept, but the lab write-up itself was pretty opaque and it felt like we were being thrown into the deep end to actually implement the networking. I spent the whole 75 minutes in lab just going over the reading and trying to figure out what exactly we were supposed to do...”

## lab difficulty

was surprised by confusion re: `recv()` function + `setTimeout()`  
oops! should have realized you haven't seen these kinds of interfaces  
before

probably need an introduction to this type of interface in lecture in  
the future

# callback-based programming (1)

```
/* library code you don't write */
void mainLoop() {
    while (true) {
        Event event = getNextEvent();
        if (event.type == RECIEVED) {
            recvd(...);
        } else if (event.type == TIMEOUT) {
            (event.timeout_function)(...);
        }
        ...
    }
}
```

## callback-based programming (2)

```
/* your code, called by library */  
void recvd(...) {  
    ...  
    setTimeout(..., timerCallback, ...);  
}  
  
void timerCallback(...) {  
    ...  
}
```



# callback-based programming

writing scripts in a webpage

many graphical user interface libraries

sometimes servers that handle lots of connections

# protocol

GET0 — start

other end acknowledges by giving data

if they don't acknowledge, you need to send again

ACK $n$

request message  $n + 1$  by acknowledging message  $n$

not quite same purpose as acknowledgments in lecture examples  
(in lab, the response is your 'acknowledgment' of your request;  
you retry if you don't get it)

## feedback re: quizzes

“I would appreciate if the quizzes were a little longer. We learn a lot in this class and I don't think 5 questions (sometimes with no partial credit) is the best representation of our skills.”

“would you ever consider dropping the lowest quiz grade?”

“I've found that the quizzes are incredibly difficult...I feel that the scope of the quizzes is way beyond the lecture material and readings...Maybe lecture material and/or readings could more closely align with the quiz questions, so that we are better prepared.”

# layers

application	HTTP, SSH, SMTP, ...	application-defined meanings
transport	TCP, UDP, ...	reach      correct      program, reliability/streams
network	IPv4, IPv6, ...	reach      correct      machine (across networks)
link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
physical	...	encode bits for wire/radio

# layers

application	HTTP, SSH, SMTP, ...	application-defined meanings
transport	TCP, UDP, ...	reach      correct      program, reliability/streams
network	IPv4, IPv6, ...	reach      correct      machine (across networks)
link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
physical	...	encode bits for wire/radio

# more than four layers?

sometimes more layers above 'application'

e.g. HTTPS:

HTTP (app layer) on TLS (another app layer) on TCP (network) on ...

e.g. DNS over HTTPS:

DNS (app layer) on HTTP on on TLS on TCP on ...

e.g. SFTP:

SFTP (app layer??) on SSH (another app layer) on TCP on ...

e.g. HTTP over OpenVPN:

HTTP on TCP on IP on OpenVPN on UDP on different IP on ...

# names and addresses

name	address
logical identifier	location/how to locate
variable counter	memory address 0x7FFF9430
DNS name www.virginia.edu	IPv4 address 128.143.22.36
DNS name mail.google.com	IPv4 address 216.58.217.69
DNS name mail.google.com	IPv6 address 2607:f8b0:4004:80b::2005
DNS name reiss-t3620.cs.virginia.edu	IPv4 address 128.143.67.91
DNS name reiss-t3620.cs.virginia.edu	MAC address 18:66:da:2e:7f:da
service name https	port number 443
service name ssh	port number 22

# layers

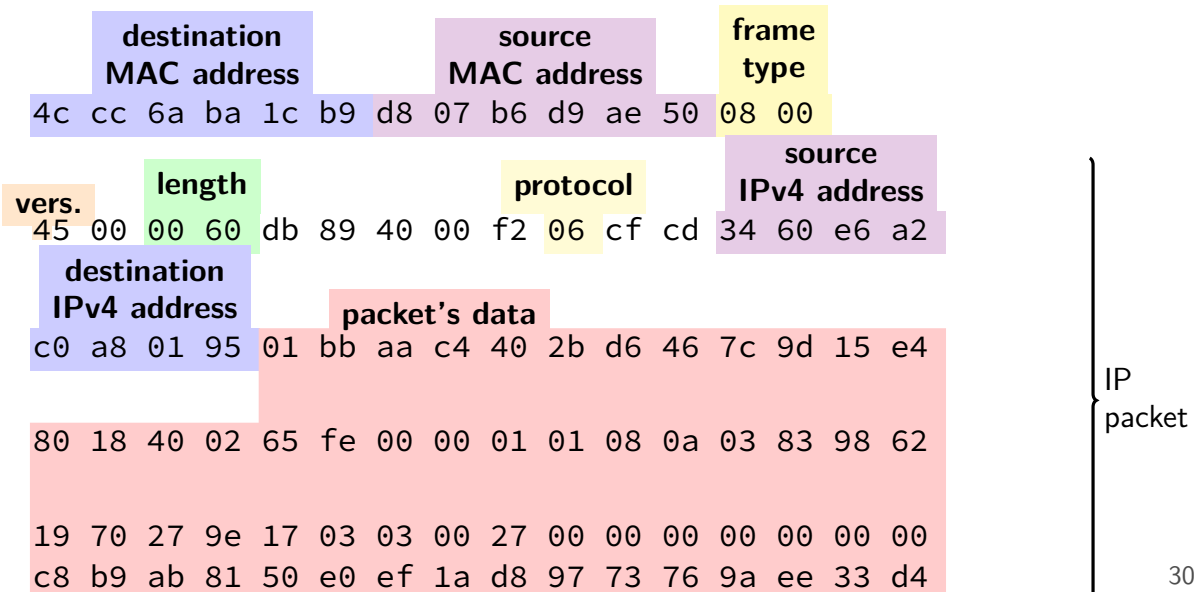
application	HTTP, SSH, SMTP, ...	application-defined meanings
transport	TCP, UDP, ...	reach      correct      program, reliability/streams
network	IPv4, IPv6, ...	reach      correct      machine (across networks)
link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
physical	...	encode bits for wire/radio



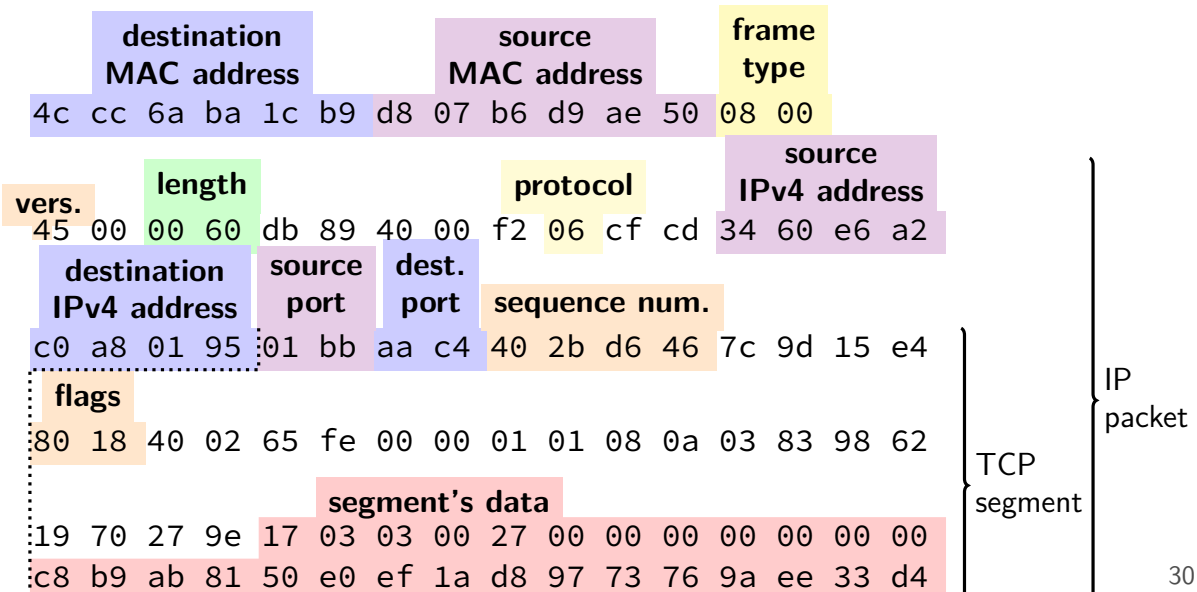
# an Ethernet frame

destination MAC address						source MAC address						frame type			
4c	cc	6a	ba	1c	b9	d8	07	b6	d9	ae	50	08	00		
frame's data															
45	00	00	60	db	89	40	00	f2	06	cf	cd	34	60	e6	a2
c0	a8	01	95	01	bb	aa	c4	40	2b	d6	46	7c	9d	15	e4
80	18	40	02	65	fe	00	00	01	01	08	0a	03	83	98	62
19	70	27	9e	17	03	03	00	27	00	00	00	00	00	00	00
c8	b9	ab	81	50	e0	ef	1a	d8	97	73	76	9a	ee	33	d4

# an Ethernet frame



# an Ethernet frame



# the link layer

Ethernet, Wi-Fi, Bluetooth, DOCSIS (cable modems), ...

allows send/recv messages to machines on “same” network segment

- typically: wireless range+channel or connected to a single switch/router
- could be larger (if *bridging* multiple network segments)
- could be smaller (switch/router uses “virtual LANs”)

typically: source+destination specified with MAC addresses

- MAC = media access control

- usually manufacturer assigned / hard-coded into device
- unique address per port/wifi transmitter/etc.

can specify destination of “anyone” (called *broadcast*)

# link layer quality of service

if frame gets...

event	on Ethernet	on WiFi
collides with another	detected + may resend	resend
not received	lose silently	resent
header corrupted	usually discard silently	usually resend
data corrupted	usually discard silently	usually resend
too long	not allowed to send	not allowed to send
reordered (v. other messages)	received out of order	received out of order
destination unknown	lose silently	usually resend??
too much being sent	discard excess?	discard excess?

# link layer reliability?

Ethernet + Wifi have checksums

Q1: Why doesn't this give us uncorrupted messages?

Why do we still have checksums at the higher layers?

Q2: What's a benefit of doing this if we're also doing it in the higher layer?

# layers

application	HTTP, SSH, SMTP, ...	application-defined meanings
transport	TCP, UDP, ...	reach      correct      program, reliability/streams
network	IPv4, IPv6, ...	reach      correct      machine (across networks)
link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
physical	...	encode bits for wire/radio

# the network layer

the Internet Protocol (IP) version 4 or version 6

there are also others, but quite uncommon today

allows send messages to/recv messages from other networks

“internetwork”

messages usually called “packets”



# network layer quality of service

if packet ...

event	on IPv4/v6
collides with another	out of scope — handled by link layer
not received	lost silently
header corrupted	usually discarded silently
data corrupted	received corrupted
too long	dropped with notice or “fragmented” + recombined
reordered (v. other messages)	received out of order
destination unknown	usually dropped with notice
too much being sent	discard excess

# network layer quality of service

if packet ...

event	on IPv4/v6
collides with another	out of scope — handled by link layer
not received	lost silently
header corrupted	usually discarded silently
data corrupted	received corrupted
too long	dropped with notice or “fragmented” + recombined
reordered (v. other messages)	received out of order
destination unknown	usually dropped with notice
too much being sent	discard excess

includes dropped by link layer  
(e.g. if detected corrupted there)

# IPv4 addresses

32-bit numbers

typically written like 128.143.67.11

four 8-bit decimal values separated by dots

first part is most significant

same as  $128 \cdot 256^3 + 143 \cdot 256^2 + 67 \cdot 256 + 11 = 2\,156\,782\,459$

organizations get blocks of IPs

e.g. UVA has 128.143.0.0–128.143.255.255

e.g. Google has 216.58.192.0–216.58.223.255 and

74.125.0.0–74.125.255.255 and 35.192.0.0–35.207.255.255

some IPs reserved for non-Internet use (127.\*, 10.\*, 192.168.\*)

# IPv6 addresses

IPv6 like IPv4, but with 128-bit numbers

written in hex, 16-bit parts, separated by colons ( : )

strings of 0s represented by double-colons ( :: )

typically given to users in blocks of  $2^{80}$  or  $2^{64}$  addresses  
no need for address translation?

`2607:f8b0:400d:c00::6a =`

`2607:f8b0:400d:0c00:0000:0000:0000:006a`


`2607f8b0400d0c000000000000000000006a`<sub>SIXTEEN</sub>

# selected special IPv6 addresses

`::1` = localhost

anything starting with `fe80` = link-local addresses  
never forwarded by routers

# IPv4 addresses and routing tables



if I receive data for...	send it to...
128.143.0.0—128.143.255.255	network 1
192.107.102.0—192.107.102.255	network 1
...	...
4.0.0.0—7.255.255.255	network 2
64.8.0.0—64.15.255.255	network 2
...	...
anything else	network 3

# selected special IPv4 addresses

127.0.0.0 — 127.255.255.255 — localhost

AKA loopback

the machine we're on

typically only 127.0.0.1 is used

192.168.0.0–192.168.255.255 and

10.0.0.0–10.255.255.255 and

172.16.0.0–172.31.255.255

“private” IP addresses

not used on the Internet

commonly connected to Internet with **network address translation**

also 100.64.0.0–100.127.255.255 (but with restrictions)

169.254.0.0–169.254.255.255

link-local addresses — ‘never’ forwarded by routers