

changelog

3 Dec 2024: add some extra exercises re: variations in cache layout and how it affects what information learned from eviction information

3 Dec 2024: add slides on EVICT+RELOAD

5 Dec 2024: “supplying own attack code?”: be more explicit that code is running partially

1-4 Dec 2025: reorder slides re: exercises; add more review of general recovery pattern; cleanup mitigation slides; generalize Spectre slides some; remove stuff re: reading from unaligned int with `intArray[x]`

check_passphrase

```
int check_passphrase(const char *versus) {  
    int i = 0;  
    while (passphrase[i] == versus[i] &&  
           passphrase[i]) {  
        i += 1;  
    }  
    return (passphrase[i] == versus[i]);  
}
```

check_passphrase

```
int check_passphrase(const char *versus) {  
    int i = 0;  
    while (passphrase[i] == versus[i] &&  
           passphrase[i]) {  
        i += 1;  
    }  
    return (passphrase[i] == versus[i]);  
}
```

number of iterations = number matching characters

leaks information about passphrase, oops!

exploiting check_passphrase (1)

guess	measured time
aaaa	100 ± 5
baaa	103 ± 4
caaa	102 ± 6
<i>daaa</i>	<i>111 ± 5</i>
eaaa	99 ± 6
faaa	101 ± 7
gaaa	104 ± 4
...	...

exploiting check_passphrase (2)

guess	measured time
daaa	102 ± 5
dbaa	99 ± 4
dcaa	104 ± 4
ddaa	100 ± 6
deaa	102 ± 4
<i>dfaa</i>	<i>109 ± 7</i>
dgaa	103 ± 4
...	...

not just timing — power analysis




Figure 19.10: Plot of the current measured during 256 single attempts to guess the first byte of a service password stored in the microcontroller at the heart of a car immobilizer (courtesy of Markus Kuhn and Sergei Skorobogatov)

not just timing — radio frequency

350 MHz, 50 MHz BW, 12 frames (160 ms) averaged




Fig. 3. Text signal received from a 440CDX laptop at 10 m distance through two intermediate offices (3 plasterboard walls). From Kuhn, "Electromagnetic

From Kuhn, "Electromagnetic Eavesdropping Risks of Flat-Panel Displays" (PET 2004)

timing and cryptography

lots of asymmetric cryptography uses big-integer math

example: multiplying 500+ bit numbers together

how do you implement that?

big integer multiplication

say we have two 64-bit integers x, y

and want to 128-bit product, but our multiply instruction only does 64-bit products

one way to multiply:

divide x, y into 32-bit parts: $x = x_1 \cdot 2^{32} + x_0$ and $y = y_1 \cdot 2^{32} + y_0$

then $xy = x_1y_12^{64} + x_1y_0 \cdot 2^{32} + x_0y_1 \cdot 2^{32} + x_0y_0$

big integer multiplication

say we have two 64-bit integers x, y

and want to 128-bit product, but our multiply instruction only does 64-bit products

one way to multiply:

divide x, y into 32-bit parts: $x = x_1 \cdot 2^{32} + x_0$ and $y = y_1 \cdot 2^{32} + y_0$

then $xy = x_1y_12^{64} + x_1y_0 \cdot 2^{32} + x_0y_1 \cdot 2^{32} + x_0y_0$

can extend this idea to arbitrarily large numbers

number of smaller multiplies depends on size of numbers!

big integers and cryptography

naive multiplication idea:

number of steps depends on size of numbers

problem: sometimes the value of the number is a secret

e.g. part of the private key

oops! revealed through timing

big integer timing attacks in practice (1)

early versions of OpenSSL (TLS implementation) had timing attack


Brumley and Boneh, "Remote Timing Attacks are Practical" (Usenix Security '03)

attacker could figure out bits of private key from timing

why? variable-time multiplication and modulus operations

got faster/slower depending on how input was related to private key
e.g. detect when attacker-value MOD private key part is close to 0

big integer timing attacks in practice (2)



(a) The zero-one gap $T_g - T_{g_{hi}}$ indicates that we can distinguish between bits that are 0 and 1 of the RSA factor q for 3 different randomly-generated keys. For clarity, bits of q that are 1 are omitted, as the x -axis can be used for reference for this case.

browsers and website leakage

web browsers run code from untrusted webpages

one goal: can't tell what other webpages you visit

some webpage leakage (1)

...as you can see [here](#), [here](#), and [here](#) ...

convenient feature 1: browser marks visited links

```
<script>
var the_color = window.getComputedStyle(
    document.querySelector('a[href=~"foo.com"]')
).color
if (the_color == ...) { ... }
</script>
```

convenient feature 2: scripts can query current color of something

some webpage leakage (1)

...as you can see [here](#), [here](#), and [here](#) ...

convenient feature 1: browser marks visited links

```
<script>
var the_color = window.getComputedStyle(
    document.querySelector('a[href=~"foo.com"]')
).color
if (the_color == ...) { ... }
</script>
```

~~convenient feature 2: scripts can query current color of something~~

fix 1: `getComputedStyle` lies about the color

fix 2: limited styling options for visited links

some webpage leakage (2)


one idea: script in webpage times loop that writes big array

variation in timing depends on *other things running on machine*

some webpage leakage (2)

one idea: script in webpage times loop that writes big array

variation in timing depends on *other things running on machine*



turns out, other webpages
create distinct “signatures”

Figure from Cook et al, “There’s Always a Bigger Fish: Clarifying Analysis of a Machine-Learning-Assisted Side-Channel Attack” (ISCA ’22)

Figure 3: Example loop-counting traces collected over 15 seconds. Darker shades indicate smaller counter values and lower instruction throughput.

side channels

observing machine operation, indirectly

common types of indirect observations

- timing

- power draw

- electromagnetic radiation

side channels can be security and privacy issues

- often used to reveal internal/secret state

our focus

we're going to focus on timing-related leaks...

and interaction with attacker input/processor design

discovery: lots of cases where features that make code/processors fast

let attackers learn much more than they should through timing

inferring cache accesses (1)

suppose I time accesses to array of chars:

reading array[0]: 3 cycles

reading array[64]: 4 cycles

reading array[128]: 4 cycles

reading array[192]: 20 cycles

reading array[256]: 4 cycles

reading array[288]: 4 cycles

...

what could cause this difference?

array[192] not in some cache, but others were

inferring cache accesses (2)

some psuedocode:

```
char array[CACHE_SIZE];  
AccessAllOf(array);  
*other_address += 1;  
TimeAccessingArray();
```

suppose during these accesses I discover that `array[128]` is slower to access

probably because `*other_address` loaded into cache + evicted it

what do we know about `other_address`? (select all that apply)

- A. same cache tag B. same cache index C. same cache offset
- D. diff. cache tag E. diff. cache index F. diff. cache offset

some complications

caches often use physical, not virtual addresses

- (and need to know about physical address to compare index bits)

- (but can infer physical addresses with measurements/asking OS)

- (often OS allocates contiguous physical addresses esp. w/‘large pages’)

storing/processing timings evicts things in the cache

- (but can compare timing with/without access of interest to check this)

processor “pre-fetching” may load things into cache before access is timed

- (but can arrange accesses to avoid triggering prefetcher and make sure to measure with memory barriers)

some L3 caches use a simple hash function to select index instead of index bits

exercise: inferring cache accesses (1)

```
char *array;  
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) {  
    *pointer += 1;  
}  
if (TimeAccessTo(&array[index]) > THRESHOLD) {  
    /* pointer accessed */  
}
```

suppose pointer is 0x1000188

and cache (of interest) is direct-mapped, 32768 (2^{15}) byte, 64-byte blocks

what array index should we check?

solution

```
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) { *pointer = 1; }  
if (TimeAccessTo(&array[index]) > THRESHOLD) { /* pointer accessed */ }
```

2^{15} byte direct mapped cache, $64 = 2^6$ byte blocks

9 index bits, 6 offset bits

0x1000188: ...0000 0001 1000 1000

array[0] starts at multiple of cache size — index 0, offset 0

to get index 6, offset 0 array[0b1 1000 0000] = array[0x180]

solution

```
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) { *pointer = 1; }  
if (TimeAccessTo(&array[index]) > THRESHOLD) { /* pointer accessed */ }
```

2^{15} byte direct mapped cache, $64 = 2^6$ byte blocks

9 index bits, 6 offset bits

0x1000188: ...0*000* *0001* *1000* 1000

array[0] starts at multiple of cache size — index 0, offset 0

to get index 6, offset 0 array[0b*1* *1000* 0000] = array[0x180]

aside

```
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) { *pointer += 1; }  
if (TimeAccessTo(&array[index]) > THRESHOLD) {  
    /* pointer accessed */  
}
```

will this detect when pointer accessed? yes

will this detect if mystery is true? not quite

...because branch prediction could started cache access

exercise: inferring cache accesses (2)

```
char *other_array = ...;
char *array;
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);
LoadIntoCache(array, CACHE_SIZE);
other_array[mystery] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (TimeAccessTo(&array[i]) > THRESHOLD) {
        /* found something interesting */
    }
}
```

other_array at 0x200400, and interesting index is $i=0x800$, then what was mystery?

solution

```
array = AllocateAlignedPhysicalMemory(CACHE_SIZE);
LoadIntoCache(array, CACHE_SIZE);
other_array[mystery] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (TimeAccessTo(&array[i]) > THRESHOLD) { ... }
}
```

at $i=0x800$: ...0000 1000 0000 0000 (cache index = 0x20)

other_array at 0x200400

Q: 0x200400 + X has cache index 0x20?

0x200400	...	0	000	0100	00	00	0000
+ X	...	?	000	0100	00	??	????
0x200400+X	...	?	000	1000	00	??	????

exercise: inferring cache accesses (2)

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) {  
    *pointer = 1;  
}  
if (TimeAccessTo(&array[index1]) > THRESHOLD ||  
    TimeAccessTo(&array[index2]) > THRESHOLD) {  
    /* pointer accessed */  
}
```

pointer is 0x1000188

cache is 2-way, 32768 (2^{15}) byte, 64-byte blocks, ??? replacement

what array indexes should we check?

reading a value — general pattern

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
other_array[mystery] += 1;  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

previous idea: learn bits of mystery that correspond to index bits
compute index bits of other_array + mystery as function of mystery
see how it matches index bits of array + i

reading a value — simpler case

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
other_array[mystery * BLOCK_SIZE] += 1;  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

if other_array is char* and starts at multiple of cache size

and array[0x6*BLOCK_SIZE] is slow to access

mystery == 0x6 + K * SET_COUNT

PRIME+PROBE

name in literature: PRIME + PROBE

PRIME: fill cache (or part of it) with values

do thing that uses cache

PROBE: access those values again and see if it's slow

(one of several ways to measure how cache is used)

coined in attacks on AES encryption

example: AES (1)

from Osvik, Shamir, and Tromer, “Cache Attacks and Countermeasures: the Case of AES” (2004)

early AES implementation used lookup tables

goal: detect index into lookup table

index depended on key + data being encrypted

tricks they did to make this work

vary data being encrypted

subtract average time to look for what changes

lots of measurements

example: AES (2)

from Osvik, Shamir, and Tromer, “Cache Attacks and Countermeasures: the Case of AES” (2004)




Fig. 5. Prime+Probe attack using 30,000 encryption calls on a 2GHz Athlon 64, attacking Linux 2.6.11 `dm-crypt`. The horizontal axis is the evicted cache set (i.e., $\langle y \rangle$ plus an offset due to the table's location) and the vertical axis is p_0 . Left: raw timings (lighter is slower). Right: after subtraction of the average timing of the cache set. The bright diagonal reveals the high nibble of $p_0 = 0x00$.

revisiting an earlier example (1)

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
LoadIntoCache(array, CACHE_SIZE);  
if (mystery) {  
    *pointer += 1;  
}  
if (TimeAccessTo(&array[index]) > THRESHOLD) {  
    /* pointer accessed */  
}
```

what if mystery is false *but* branch mispredicted?

revisiting an earlier example (2)

	cycle #	0	1	2	3	4	5	6	7	8	9	10	11
movq mystery, %rax		F	D	R	I	E	E	E	W	C			
test %rax, %rax		F	D	R				I	E	W	C		
jz skip (mispred.)		F	D	R				I	E	W	C		
mov pointer, %rax		F	D	R	I	E	E	E	W				
mov (%rax), %r8			F	D	R				I	E	W		
add \$1, %r8			F	D	R								
mov %r8, %rax				F	D	R							
...													
skip: ...								F	D	R			

avoiding/triggering this problem

```
if (something false) {  
    access *pointer;  
}
```

what can we do to make access more/less likely to happen?

reading a value without really reading it

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
if (something false) {  
    other_array[mystery * BLOCK_SIZE] += 1;  
}  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

if branch mispredicted, cache access may *still happen*

can find the value of mystery

seeing past a segfault? (1)

```
Prime();  
if (something false) {  
    triggerSegfault();  
    Use(*pointer);  
}  
Probe();
```

could cache access for `*pointer` still happen?

yes, if:

- branch for if statement mispredicted, and
- `*pointer` starts before segfault detected

seeing past a segfault? (2)

operations in virtual memory lookup:

- translate virtual to physical address

- check if access is permitted by permission bits

Intel processors: looks like these were separate steps, so...

```
Prime();  
if (something false) {  
    int value = ReadMemoryMarkedNonReadableInPageTable();  
    access other_array[value * ...];  
}  
Probe();
```

seeing past a segfault? (2)

operations in virtual memory lookup:

- translate virtual to physical address

- check if access is permitted by permission bits

Intel processors: looks like these were separate steps, so...

```
Prime();  
if (something false) {  
    int value = ReadMemoryMarkedNonReadableInPageTable();  
    access other_array[value * ...];  
}  
Probe();
```

seeing past a segfault? (2)

operations in virtual memory lookup:

- translate virtual to physical address

- check if access is permitted by permission bits

Intel processors: looks like these were separate steps, so...

```
Prime();  
if (something false) {  
    int value = ReadMemoryMarkedNonReadableInPageTable();  
    access other_array[value * ...];  
}  
Probe();
```

seeing past a segfault? (2)

operations in virtual memory lookup:

- translate virtual to physical address

- check if access is permitted by permission bits

Intel processors: looks like these were separate steps, so...

```
Prime();  
if (something false) {  
    int value = ReadMemoryMarkedNonReadableInPageTable();  
    access other_array[value * ...];  
}  
Probe();
```

Meltdown

from Lipp et al, "Meltdown: Reading Kernel Memory from User Space"

```
// %rcx = kernel address  
// %rbx = array to load from to cause eviction  
xor %rax, %rax      // rax ← 0
```

retry:

```
// rax ← memory[kernel address] (segfaults)  
// but check for segfault done out-of-order on Intel at t  
movb (%rcx), %al  
// rax ← memory[kernel address] * 4096 [speculated]  
shl $0xC, %rax  
jz retry            // not-taken branch  
// access array[memory[kernel address] * 4096]  
mov (%rbx, %rax), %rbx
```

Meltdown

from Lipp et al, "Meltdown: Reading Kernel Memory from User Space"

```
// %rcx = kernel address
// %rbx = array base address
xor %rax, %rax
retry:
// rax ← memory[kernel address] (segfaults)
// but check for segfault done out-of-order on Intel at t
movb (%rcx), %al
// rax ← memory[kernel address] * 4096 [speculated]
shl $0xC, %rax
jz retry // not-taken branch
// access array[memory[kernel address] * 4096]
mov (%rbx, %rax), %rbx
```

space out accesses by 4096
ensure separate cache sets and
avoid triggering prefetcher

Meltdown

from Lipp et al, "Meltdown: Reading Kernel Memory from User Space"

```
// %rcx
// %rbx
xor %rax, %rax
retry:
    // rax <- memory[kernel address] (segfaults)
    // but check for segfault done out-of-order on Intel at t
    movb (%rcx), %al
    // rax <- memory[kernel address] * 4096 [speculated]
    shl $0xC, %rax
    jz retry
    // access array[memory[kernel address] * 4096]
    mov (%rbx, %rax), %rbx
```

repeat access if zero
apparently value of zero speculatively read
when real value not yet available

Meltdown

from Lipp et al, "Meltdown: Reading Kernel Memory from User Space"

```
// %rcx : access cache to allow measurement later  
// %rbx : in paper with FLUSH+RELOAD instead of PRIME+PROBE technique  
xor %rax, %rax  
retry:  
// rax <- memory[kernel address] (segfaults)  
// but check for segfault done out-of-order on Intel at t  
movb (%rcx), %al  
// rax <- memory[kernel address] * 4096 [speculated]  
shl $0xC, %rax  
jz retry // not-taken branch  
// access array[memory[kernel address] * 4096]  
mov (%rbx, %rax), %rbx
```

Meltdown

from Lipp et al

```
// %rax
```

```
// %rax
```

```
xor %rax, %rax
```

retry:

```
// rax ← memory[kernel address] (segfaults)
```

```
// but check for segfault done out-of-order on Intel at t
```

```
movb (%rcx), %al
```

```
// rax ← memory[kernel address] * 4096 [speculated]
```

```
shl $0xC, %rax
```

```
jz retry // not-taken branch
```

```
// access array[memory[kernel address] * 4096]
```

```
mov (%rbx, %rax), %rbx
```

segfault actually happens eventually

option 1: okay, just start a new process every time
(don't even need to hide segfault)

option 2: way of suppressing segfault

(paper used transactional memory support,
not mispredicted branch)

Meltdown fix

HW: permissions check done with/before physical address lookup
was already done by AMD, ARM apparently?
now done by Intel

SW: separate page tables for kernel and user space
don't have sensitive kernel memory pointed to by page table
when user-mode code running
unfortunate performance problem
exceptions start with code that switches page tables

Spectre

Meltdown: address translation without permissions check

seems relatively easy to fix in hardware

but idea of leaks from speculative execution has much harder to fix attacks...

reading a value without really reading it

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
if (something false) {  
    other_array[mystery * BLOCK_SIZE] += 1;  
}  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

if branch mispredicted, cache access may *still happen*

can find the value of mystery

mistraining branch predictor?

```
if (something) {  
    CodeToRunSpeculatively()  
}
```

useful for attacks:

have 'something' be false, but predicted as true

one way: run lots of times with something true then do actually run with something false

another way: learn how branch prediction caches work, run code that fills in caches in known way

reading a value — general pattern

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
other_array[mystery] += 1;  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

previous idea: learn bits of mystery that correspond to index bits
compute index bits of other_array + mystery as function of mystery
see how it matches index bits of array + i

reading a value — simpler case

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
other_array[mystery * BLOCK_SIZE] += 1;  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

if other_array is char* and starts at multiple of cache size

and array[0x6*BLOCK_SIZE] is slow to access

mystery == 0x6 + K * SET_COUNT

reading a value — * X * BLOCK_SIZE

```
char *array;  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
other_array[mystery * 8 * BLOCK_SIZE] += 1;  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

if other_array is char* and starts at multiple of cache size

and array[0x100*BLOCK_SIZE] is slow to access

mystery * 8 = 0x100 + K * SET_COUNT

mystery = 32 + K * SET_COUNT / 8

contrived(?) vulnerable code (1)

suppose this C code is run with extra privileges

(e.g. in system call handler, library called from JavaScript in webpage, etc.)

assume x chosen by attacker

(example from original Spectre paper)

```
if (x < array1_size)
    y = array2[array1[x] * 4096];
```

the out-of-bounds access (1)

```
char array1[...];
```

```
...
```

```
int secret;
```

```
...
```

```
y = array2[array1[x] * 4096];
```

suppose array1 is at 0x10000000 and

secret is at 0x103F0003;

what x do we choose to make array1[x] access first byte of secret?

the out-of-bounds access (2)

```
unsigned char array1[...];
```

```
...
```

```
int secret;
```

```
...
```

```
y = array2[array1[x] * 4096];
```

suppose our cache has 64-byte blocks and 8192 sets

and `array2[0]` is char stored in cache set 0

if the above evicts something in cache set 128,
then what do we know about `array1[x]`?

the out-of-bounds access (2)

```
unsigned char array1[...];
```

```
...
```

```
int secret;
```

```
...
```

```
y = array2[array1[x] * 4096];
```

suppose our cache has 64-byte blocks and 8192 sets

and `array2[0]` is char stored in cache set 0

if the above evicts something in cache set 128,
then what do we know about `array1[x]`?

is 2 or 130

exploit with contrived(?) code

```
/* in kernel: */  
int systemCallHandler(int x) {  
    if (x < array1_size)  
        y = array2[array1[x] * 4096];  
    return y;  
}
```

```
/* exploiting pseudocode */  
    /* step 1: mistrain branch predictor */  
for (a lot) {  
    systemCallHandler(0 /* less than array1_size */);  
}  
  
    /* step 2: evict from cache using misprediction */  
Prime();  
systemCallHandler((targetAddress - array1Address) / A1ElemSize);  
int evictedSet = ProbeAndFindEviction();  
int targetValue = (evictedSet - array2StartSet) / setsPer4KA2Elem;
```

really contrived?

```
char *array1; char *array2;  
if (x < array1_size)  
    y = array2[array1[x] * 4096];
```

times 4096 shifts so we can get lower bits of target value
so all bits effect what cache block is used

really contrived?

```
char *array1; char *array2;  
if (x < array1_size)  
    y = array2[array1[x] * 4096];
```

times 4096 shifts so we can get lower bits of target value
so all bits effect what cache block is used

```
int *array1; int *array2;  
if (x < array1_size)  
    y = array2[array1[x]];
```

will still get *upper* bits of array1[x] (can tell from cache set)
still likely to be sensitive data

bounds check in kernel

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

bounds check in kernel

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

bounds check in kernel

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

bounds check in kernel

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

generalizing exploit

limited in what address we can learn about based on how big entries in tables are

- but can combine multiple Spectre-type exploits
- only need one secret value leaked

need to adjust calculations to actual addresses / array element sizes / etc.

privilege levels?

vulnerable code runs with higher privileges

so far: higher privileges = kernel mode

but other common cases of higher privileges

example: scripts in web browsers

JavaScript

JavaScript: scripts in webpages

not supposed to be able to read arbitrary memory, but...

can access arrays to examine caches

and could take advantage of some browser function being vulnerable

JavaScript

JavaScript: scripts in webpages

not supposed to be able to read arbitrary memory, but...

can access arrays to examine caches

and could take advantage of some browser function being vulnerable

or — *doesn't even need browser to supply vulnerable code itself!*

just-in-time compilation?

for performance, compiled to machine code, run in browser

not supposed to be access arbitrary browser memory

example JavaScript code from paper:

```
if (index < simpleByteArray.length) {  
    index = simpleByteArray[index | 0];  
    index = (((index * 4096) | 0) & (32*1024*1024-1)) | 0;  
    localJunk ^= probeTable[index|0] | 0;  
}
```

web page runs a lot to train branch predictor

then does run with out-of-bounds index

examines what's evicted by probeTable access

supplying own attack code?

JavaScript: could supply own attack code

turns out also possible with kernel mode scenario

trick: don't need to *actually run* code for real

...just need branch predictor to fetch it
so it gets partially executed speculatively

other misprediction

so far: talking about mispredicting direction of branch

what about mispredicting target of branch in, e.g.:

```
// possibly from C code like:  
//      (*function_pointer)();  
jmp *%rax
```

```
// possibly from C code like:  
//      switch(rcx) { ... }  
jmp *(%rax,%rcx,8)
```

an idea for predicting indirect jumps

for jmps like `jmp *%rax` predict target with cache:

bottom 12 bits of jmp address	last seen target
-------------------------------	------------------

0x0–0x7	0x200000
---------	----------

0x8–0xF	0x440004
---------	----------

0x10–0x18	0x4CD894
-----------	----------

0x18–0x20	0x510194
-----------	----------

0x20–0x28	0x4FF194
-----------	----------

...

...

0xFF8–0xFFF	0x3F8403
-------------	----------

Intel Haswell CPU did something similar to this

uses Hash(bits of last several jumps), not just last jmp

using mispredicted jump

- 1: find some kernel function with `jmp *%rax`
- 2: mistrain branch target predictor for it to jump to chosen code
 - use code at address that conflicts in “recent jumps cache”
 - since only bottom bits are used, can set this up in user memory
- 3: have chosen code be attack code (e.g. array access)
 - either write special code OR
 - find suitable instructions (e.g. array access) in existing kernel code
- 4: run the kernel function

Spectre variants

showed Spectre variant 1 (array bounds), 2 (indirect jump)
from original paper

other possible variations:

- could cause other things to be mispredicted

 - prediction of where functions return to?

 - values instead of which code is executed?

- could use side-channel other than data cache changes

 - instruction cache

 - cache of pending stores not yet committed

 - contention for resources on multi-threaded CPU core

 - branch prediction changes

 - ...

some Linux kernel mitigations (1)

replace `array[x]` with
`array[x & ComputeMask(x, size)]`

...where `ComputeMask()` returns

0 if $x > \text{size}$

`0xFFFF..F` if $x \leq \text{size}$

...and `ComputeMask()` does not use jumps:

```
mov x, %r8
mov size, %r9
cmp %r9, %r8
sbb %rax, %rax // sbb = subtract with borrow
                // either 0 or -1
```

some Linux kernel mitigations (2)

for indirect branches:

with hardware help:

- separate indirect (computed) branch prediction for kernel v user mode
- other branch predictor changes to isolate better

without hardware help:

- transform `jmp *(%rax)`, etc. into code that will only be predicted to jump to safe locations (by writing assembly very carefully)

only safe prediction

as replacement for `jmp *%rax`

code from Intel's "Retpoline: A Branch Target Injection Mitigation"

```
    call load_label
capture_ret_spec:    /* <-- want prediction to go here */
    pause
    lfence
    jmp capture_ret_spec
load_label:
    mov %rax, (%rsp)
    ret
```

not just BLOCK_SIZE

```
char *array, *other_array;  
// PRIME  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
// (some code we don't control)  
other_array[mystery * N] += 1; // previously: * BLOCK_SIZE  
// PROBE  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

64KB (2^{16} B) direct-mapped cache with 64B blocks

array[0x800] slow to access?

other_array at 0x40000000 (index 0, offset 0)

value of mystery if $N = 1$? $N = 32 * 64$?

solution (N=1)

$$\lfloor \text{mystery} * N / \text{BLOCK_SIZE} \rfloor \bmod 1024 = 32$$

$$\lfloor \text{mystery} * N / \text{BLOCK_SIZE} \rfloor = 32 + 1024K$$

let offset be some number in $[0, \text{BLOCK_SIZE})$:

$$\text{mystery} * N = \text{BLOCK_SIZE} \times (32 + 1024Z) + \text{offset}$$

$$\text{mystery} = \text{BLOCK_SIZE} \times (32 + 1024Z) + N \times \text{offset}$$

$$\text{mystery} = 64 \times (32 + 1024Z) + N \times \text{offset}$$

N=1: mystery = 2048, 2049, 2050, ..., 2048 + 63, $64 \cdot 1024 + 2048$,
 $64 \cdot 1024 + 2048 + 1$, ...

exercise (N=32*64)

what if $N = 32 \cdot 64$

recall: `other_array[0]` is set 0, offset 0

`other_array[mystery * N]` is set 32

possible values of mystery?

$$\text{mystery} \cdot 32 \cdot 64 = 64(32 + 1024Z) + \text{offset}$$

$$= 64 \cdot 32 + 65536Z + \text{offset}$$

$$\text{mystery} = 1 + \frac{65536}{64 \cdot 32}Z + \frac{\text{offset}}{64 \cdot 32} = 1 + 32Z$$

alternate view

learn index bits of mystery * N

this example: bits 6–15

$N = 1$, bits 6–15 of mystery

$N = 64$, bits 0–9 of mystery

$N = 32 * 64 (2^{11})$, bits 0–4 of mystery

exercise

```
char *array;  
// PRIME  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
// (some code we don't control)  
other_array[mystery * BLOCK_SIZE] += 1;  
// PROBE  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) {  
        ...  
    }  
}
```

64KB (2^{16} B) direct-mapped cache with 64B blocks

array[0x800] slow to access;

other_array at 0x40000000

value of mystery?

exercise solution (1)

$\text{NUM_SETS} = 64\text{KB}/64\text{B} = 1\text{K} (1024) \text{ sets}$

$\text{array}[0\text{x}800]$ has cache index $0\text{x}800/\text{BLOCK_SIZE} \bmod \text{NUM_SETS}$
 $= \text{cache index } 32$

know $\text{other_array}[\text{mystery} * \text{BLOCK_SIZE}]$ had same index

$\text{other_array}[0]$ at cache index 0
 $(0\text{x}4000000 / \text{BLOCK_SIZE}) \bmod \text{NUM_SETS} = 0$

exercise solution (2)

recall have found:

`other_array[0]` at index 0;

`other_array[mystery*BLOCK_SIZE]` has index 32 (same as `array[0x800]`)

`other_array[X]` at cache index $(0 + X/\text{BLOCK_SIZE} \bmod \text{NUM_SETS})$

advanced by $X/\text{BLOCK_SIZE}$ blocks

wrapping around after `NUM_SETS` blocks

$$X = \text{mystery} * \text{BLOCK_SIZE}$$

$$32 = 0 + \text{mystery} \bmod \text{NUM_SETS}$$

$$\text{mystery} = 32 \text{ or } 32 \pm 1024 \text{ or } 32 \pm 1024 \times 2 \text{ or etc.}$$

exercise

```
char *array;
//PRIME
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array);
other_array[mystery] += 1;
//PROBE
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) {
        ...
    }
}
```

with 64KB direct-mapped cache with 64B blocks

suppose we find out that `array[0x200]` is slow to access

and `other_array` starts at some multiple of cache size

What was mystery?

```

char *array;
//PRIME
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array); // PRIME
other_array[mystery] += 1;
//PROBE
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) // PROBE
        {...}
}

```

- $NSETS = CACHE_SIZE / BLOCK_SIZE = 64KB / 64B = 1K = 2^{10}$
- And this affected `array[0x200]`
 - Which had cache index $0x200 / BLOCK_SIZE = 512 / 64 = 8$
 - Or `0b 0010 0000 0000`
- `other_array[mystery] = other_array + mystery` (because these are char array)
- If we know the base address of `other_array` is `0x20000`, we need to index $(0x20000 + mystery) = 8$
- `0b 0010 0000 0000 0000 0000 //other_array`
- `+0b 0000 0000 0000 0000 //mystery`
- `=0b 0000 0000 0000 0000`
- So we get a couple bits in the low-order byte of `mystery` and the next byte

variation: different starting location

other_array starts at 0x4001440

then other_array[0] at cache index

$$0x4001440 / \text{BLOCK_SIZE} \bmod \text{NUM_SETS} = 0x51$$

$$(0x51 + \text{mystery} * \text{BLOCK_SIZE} / \text{BLOCK_SIZE}) \bmod \text{NUM_SETS} = 32$$

mystery = -49 or 975 or 1099 or ...

variation: associative cache

```
char *array;  
// PRIME  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
// (some code we don't control)  
other_array[mystery * BLOCK_SIZE] += 1;  
// PROBE  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) { ... }  
}
```

suppose 2-way 64KB cache instead of direct-mapped

$\text{NUM_SETS} = 64\text{KB} / 2 / 64\text{B} = 512$ sets

array[0x800] still has cache index 32 (still)

but now mystery can be 32 or $32 + 512$ or $32 + 512 \cdot 2$ or ...

variation: associative cache (2)

```
char *array;  
// PRIME  
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);  
AccessAllOf(array);  
// (some code we don't control)  
other_array[mystery * BLOCK_SIZE] += 1;  
// PROBE  
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {  
    if (CheckIfSlowToAccess(&array[i])) { ... }  
}
```

suppose 2-way 64KB cache w/ 64B and `array[0x8800]` is slow

$$0x8800 / \text{BLOCK_SIZE} = 544 = 512 + 32$$

since 512 sets total, still set index 32

mystery still 32 or $32 + 512$ or $32 + 512 \cdot 2$ or ...

exercise

if 4-way 64KB cache w/64B blocks and something from cache set 32 evicted,
then where could slow access be?

recall: 2-way cache: $i=0x800$, $i=0x8800$

- A. $i=0x400$, $i=0x800$, $i=0x8400$, $i=0x8800$
- B. $i=0x800$, $i=0x8800$, $i=0x10800$, $i=0x18800$
- C. $i=0x800$, $i=0x4800$, $i=0x8800$, $i=0xc800$
- D. $i=0x800$, $i=0x4800$, $i=0x8800$, $i=0x10800$
- E. something else

EVICT+RELOAD

PRIME+PROBE: fill cache, detect eviction

alternate idea EVICT+RELOAD:

```
unsigned char *probe_array;  
posix_memalign(&probe_array, CACHE_SIZE, CACHE_SIZE);  
access OTHER things to evict all of probe_array  
if (something false) {  
    read probe_array[mystery * BLOCK_SIZE];  
}
```

check which value from probe_array is faster

requires code to access something you can access

but often easier to setup/more reliable than PRIME+PROBE

EVICT+RELOAD

PRIME+PROBE: fill cache, detect eviction

alternate idea *EVICT*+RELOAD:

```
unsigned char *probe_array;  
posix_memalign(&probe_array, CACHE_SIZE, CACHE_SIZE);  
access OTHER things to evict all of probe_array  
if (something false) {  
    read probe_array[mystery * BLOCK_SIZE];  
}
```

check which value from probe_array is faster

requires code to access something you can access

but often easier to setup/more reliable than PRIME+PROBE

EVICT+RELOAD

PRIME+PROBE: fill cache, detect eviction

alternate idea EVICT+*RELOAD*:

```
unsigned char *probe_array;  
posix_memalign(&probe_array, CACHE_SIZE, CACHE_SIZE);  
access OTHER things to evict all of probe_array  
if (something false) {  
    read probe_array[mystery * BLOCK_SIZE];  
}
```

check which value from probe_array is faster

requires code to access something you can access

but often easier to setup/more reliable than PRIME+PROBE

into exploit: Meltdown

```
uint8_t* probe_array = new uint8_t[256 * 4096];  
// ... Make sure probe_array is not cached  
uint8_t kernel_memory_val = *(uint8_t*)(kernel_address);  
uint64_t final_kernel_memory = kernel_memory_val * 4096;  
uint8_t dummy = probe_array[final_kernel_memory];  
// ... catch page fault  
// ... in signal handler, determine which of 256 slots in probe_array
```

backup slides

Quiz week 14

Question 5 (0 / 4 pt; mean 0.16)

Consider the following C code:

```
struct Files all_files[NUM_FILES];
int GetLastByteOfFile(int index) {
    if (index >= 0 && index < all_files) {
        struct File *file = &all_files[index];
        if (file->type == MEMORY) {
            return file->data[file->size - 1];
        } else if (file->type == DISK) {
            return GetLastByteOfDiskFile(file);
        } else {
            return -1;
        }
    } else {
        return -1;
    }
}
```

If the above function runs in kernel mode, we might be able to use a Spectre-style attack where the cache evictions caused by memory access of `file->data[file.size - 1]` allows us learn about the value of an arbitrary memory location. To perform this attack, the attacker would prefer to choose an out-of-bounds index such that ____.

- A. 88% ☐ the address of `all_files[index].data[file.size - 1]` is the memory address whose value they want to learn about
- B. 2% ☐ the address of `all_files[index].type` is the memory address they want to learn about
- C. 4% ^T
(correct) ☒ the address of `all_files[index].size` is the memory address they want to learn about
- D. ☐ the value of `all_files[index].type` is DISK

- `file = &all_files[index]`
- So if we provide an out of bounds index, we can read arbitrary memory
- `file->data[]`, ie `all_files[index].data[]` is like array2
- `file->size`, ie `all_files[index].size`, is like array1

```
if (x < array1_size) {  
    y = array2[array1[x]];  
}
```

our template

```
void SomeSystemCallHandler(int index) {  
    if (index > some_table_size)  
        return ERROR;  
    int kind = table[index];  
    switch (other_table[kind].foo) {  
        ...  
    }  
}
```

actual code

- kind ~ file->size, ie all_files[index].size, is like array1
 - This is the address we want to learn about by observing its cache behavior
- other_table [] ~ file->data[], ie all_files[index].data[] is like array2
 - Not using the .foo here

Question 5 (0 / 4 pt; mean 0.16)

Consider the following C code:

```
struct Files all_files[NUM_FILES];
int GetLastByteOfFile(int index) {
    if (index >= 0 && index < all_files) {
        struct File *file = &all_files[index];
        if (file->type == MEMORY) {
            return file->data[file->size - 1];
        } else if (file->type == DISK) {
            return GetLastByteOfDiskFile(file);
        } else {
            return -1;
        }
    } else {
        return -1;
    }
}
```

If the above function runs in kernel mode, we might be able to use a Spectre-style attack where the cache evictions caused by memory access of `file->data[file.size - 1]` allows us learn about the value of an arbitrary memory location. To perform this attack, the attacker would prefer to choose an out-of-bounds index such that ____.

- A. 88% ☐ the address of `all_files[index].data[file.size - 1]` is the memory address whose value they want to learn about
- B. 2% ☐ the address of `all_files[index].type` is the memory address they want to learn about
- C. 4% ^T
(correct) ☒ the address of `all_files[index].size` is the memory address they want to learn about
- D. ☐ the value of `all_files[index].type` is DISK

- Why not A?
- `all_files[index].data` is like `array2`
 - Based on which cache set is affected by different index values, we learn what those index values are
 - So we need to set up the index (`~array1`) to refer to the memory location we're interested in – here `file.size`, ie `all_files[index].size`

Q4

Consider the following code:

```
unsigned char check_array[32768];
int mystery = /* unknown */;

int Check(int key) {
    return check_array[(key + mystery) % 32768];
}
```

Suppose that:

- (to simplify the problem) virtual memory is not use
- check_array is located at physical address 0x1400000
- the system has a 2-way 64KB (2 to 16 byte) data cache with 64-byte cache blocks.
- Check is compiled to perform exactly three memory accesses:
 - to read the global variable read mystery
 - to reads its return address from the stack
 - to read from check_array

Suppose we determine that calling Check evicts from the data cache sets as follows:

key value | evicts values from cache set indexes

0	15, 72, 435
32	15, 72, 435
48	15, 72, 436
128	15, 72, 437
8240	15, 52, 72

Based on this information what is a possible value for mystery?

There was a typo originally that we fixed.
Originally wrote %16384 instead of %32768,
and 4096 instead of 8240 for the last row

- $(0 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$
- $(32 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$
- $(48 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 436$
- ...

Suppose we determine that calling Check evicts from the data cache sets as follows:

key value	evicts values from cache set indexes
0	15, 72, 435
32	15, 72, 435
48	15, 72, 436
128	15, 72, 437
8240	15, 52, 72

Based on this information what is a possible value for mystery?

Answer:

Key: a value between 27856 and 27871 +/- any multiple of 32768

originally we erroneously wrote % 16384 instead of % 32768, and 4096 instead of 8240 for the last value

$(0 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$

$(32 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 435$

$(48 + \text{mystery}) // 64 \text{ +/- multiple of } 512 \text{ (number of cache sets)} = 436$

...

let's choose the multiple of 512 to be 0 for simplicity, for now

$\text{mystery} // 64 = 435$ implies mystery in $[435 * 64 = 27840, 27840 + 63 = 27903]$

$(32 + \text{mystery}) // 64 = 435$ implies mystery in $[435 * 64 - 32 = 27808, 27808 + 63 = 27871]$

$(48 + \text{mystery}) // 64 = 436$ implies mystery in $[436 * 64 - 48 = 27856, 27856 + 63 = 27919]$

overlap here implies mystery in $[27856, 27871]$

the multiple of 512 offsets this by $512 * 64 = 32768$

extracting low-order bits

```
char *array;
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array);
other_array[mystery * BLOCK_SIZE] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) {
        ...
    }
}
```

with 64KB direct-mapped cache with 64B blocks

suppose we find out that `array[0x700]` is slow to access

and `other_array` starts at some multiple of cache size

What was mystery?

```

char *array;
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array); // PRIME
other_array[mystery] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) // PROBE
        {...}
}

```

- $NSETS = CACHE_SIZE / BLOCK_SIZE = 64KB / 64B = 1K = 2^{10}$
- And this affected `array[0x700]` //cache-aligned
 - Which had cache index $0x700 / BLOCK_SIZE = 1792 / 64 = 28$
 - Or `0b 0111 0000 0000`
- `other_array[mystery] = other_array + mystery` (because these are char array)
- If we know the base address of `other_array` is `0x20000`, we need `index(0x20000 + mystery) = 28`
- `0b 0010 0000 0000 0000 0000` //other_array
- `+0b 0000 0000 0000 0000 0000` //mystery
- `=0b 0010 0000 0000 0000 0000`
- Now we find the low order byte of mystery, which is `0b 0001 1100 = 28`
- In either case, we extract $\log(NSETS)$ bits, at the positions that align with the index bits

```

char *array;
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array); // PRIME
other_array[mystery] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) // PROBE
        {...}
}

```

- $NSETS = CACHE_SIZE / BLOCK_SIZE = 64KB / 64B = 1K = 2^{10}$
- And this affected `array[0x700]`
 - Which had cache index $0x700 / BLOCK_SIZE = 1792 / 64 = 28$
 - Or `0b 0111 0000 0000`
- `other_array[mystery] = other_array + mystery` (because these are char array)
- If we know the base address of `other_array` is `0x20440`, we need to index $(0x20440 + mystery) = 28$
- `0b 0010 0000 0100 0100 0000` //other_array
- `+0b 0000 0000 0010 11??` //mystery
- `=0b 0010 0000 0110 11??`
- Now we find the actual value of mystery, which is `0b 0000 1011 = 11`

```

char *array;
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array); // PRIME
other_array[mystery * BLOCK_SIZE] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) // PROBE
        {...}
}

```

- $NSETS = CACHE_SIZE / BLOCK_SIZE = 64KB / 64B = 1K$
- Each value of `mystery` touches a different cache line
 - So we touched cache index $mystery \% NSETS$
 - But base address might be offset
- And this affected `array[0x700]`
 - Which had cache index $0x700 / BLOCK_SIZE = 1792 / 64 = 28$
- And `&other_array` starts at 0x20440, which has cache index $(0x20440 / BLOCK_SIZE) \% NSETS = 17$
- So $IDX(mystery) + IDX(\&other_array) = 28$
- So $IDX(mystery) = 28 - 17 = 11$
- So `mystery = 11` or $(11 + 1024)$ or ...
 - If we know `mystery` is a char, then we know it's between 0-255, so in this case `mystery = 11`
- It's the same math!!!

```

char array[CACHE_SIZE] // not aligned
AccessAllOf(array); // PRIME
other_array[mystery * BLOCK_SIZE] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) // PROBE
        {...}
}

```

- $NSETS = CACHE_SIZE / BLOCK_SIZE = 64KB / 64B = 1K$
- Each value of `mystery` touches a different cache line
 - So we touched cache index `mystery % NSETS`
 - But base address might be offset
- And this affected `array[0x8280]`
 - Whose base address might also be offset, say `0x48480`
 - What cache index is `array[0x8280]`?
 - $IDX(\&array + 0x8280) = ((0x48480 + 0x8280) / BLOCK_SIZE) \% NSETS = 28$
- And `&other_array` starts at `0x20440`, which has cache index $(0x20440 / BLOCK_SIZE) \% NSETS = 17$
- So $IDX(mystery) + IDX(\&other_array) = 28$
- So $IDX(mystery) = 28 - 17 = 11$
- So `mystery = 11` or $(11 + 1024)$ or ...
 - If we know `mystery` is a char, then we know it's between 0-255, so in this case `mystery = 11`

What about associative caches?

```
char *array;
posix_memalign(&array, CACHE_SIZE, CACHE_SIZE);
AccessAllOf(array);
other_array[mystery * BLOCK_SIZE] += 1;
for (int i = 0; i < CACHE_SIZE; i += BLOCK_SIZE) {
    if (CheckIfSlowToAccess(&array[i])) {
        ...
    }
}
```

with 64KB 2-way cache with 64B blocks

suppose we find out that `array[0x800]` is slow to access

and `other_array` starts at some multiple of cache size

What was mystery?

another exercise

```
char array1[...];  
...  
int secret;  
...  
y = array2[array1[x] * 4096];
```

- Suppose our cache has 64B blocks and 1K sets, and array2[0] is in set 0
- Suppose our prime+probe lets us see that something in cache set 256 or our probe array (array2) is evicted
- What do we know about array1[x]?

```
char array1[...];  
...  
int secret;  
...  
y = array2[array1[x] * 4096];
```

- Suppose our cache has 64B blocks and 1K sets, and array2[0] is in set 0
 - So array2[64] is in set 1, array2[128] is in set 2, etc.
- Suppose our prime+probe lets us see that something in cache set 256 of our probe array (array2) is evicted,
 - So $\text{CACHE_SET}(\text{array1}[x] * 4096) = 256$
- What do we know about array1[x]?
 - $\text{array1}[x] * 4K = 64 * \text{target_set} + \text{some multiple of number of sets}$
 - $\text{array1}[x] * 4K = 64 * 256 + \dots$
 - So $\text{array1}[x] = (64 * 256) / 4K = 16K / 4K = 4 + \dots$

```
char array1[...];  
...  
int secret;  
...  
y = array2[array1[x] * 4096];
```

- Suppose our cache has 64B blocks and 32K sets, and array2[0] is in set 0
 - So array2[64] is in set 1, array2[128] is in set 2, etc.
- Suppose our prime+probe lets us see that something in cache set 256 of our probe array is evicted, so $\text{CACHE_SET}(\text{array1}[x] * 4096) = 256$
- What do we know about array1[x]?
 - $\text{array1}[x] * 4K = 64 * \text{target_set} + \text{some multiple of number of sets}$
 - $\text{array1}[x] * 4K = 64 * 256 + n * 32K * 64$
 - So $\text{array1}[x] = (64 * 256 + n * 32K * 64) / 4K = 16K / 4K + (n * 32K * 64) / 4K$
 - So array1[x] = 4 or 4+512 or...
 - But it's a char, so it can only be 4

```
char array1[...];  
...  
int secret;  
...  
y = array2[array1[x] * 4096];
```

- Suppose our cache has 64B blocks and 2K sets, and array2[0] is in set 0
 - So array2[64] is in set 1, array2[128] is in set 2, etc.
- Suppose our prime+probe lets us see that something in cache set 256 of our probe array is evicted, so $\text{CACHE_SET}(\text{array1}[x] * 4096) = 256$
- What do we know about array1[x]?
 - $\text{array1}[x] * 4K = 64 * \text{target_set} + \text{some multiple of number of sets}$
 - $\text{array1}[x] * 4K = 64 * 256 + n * 2K * 64$
 - So $\text{array1}[x] = (64 * 256 + n * 2K * 64) / 4K = 16K / 4K + (n * 2K * 64) / 4K$
 - So array1[x] = 4 or 4+32 or 4+64 or...
 - But it's a char, so it can only be 4, 36, 68, 100, 132, 164, or 196
 - ... This works better in last-level caches with larger # of sets

```

char array1[...];
...
int secret;
...
y = array2[array1[x]]; // no *4096 this time

```

- Suppose our cache has 64B blocks and 32K sets, and array2[0] is in set 0
 - So array2[64] is in set 1, array2[128] is in set 2, etc.
- Suppose our prime+probe lets us see that something in cache set **3** of our probe array is evicted, so $\text{CACHE_SET}(\text{array1}[x] \text{ ~~*4096~~}) = 3$
- What do we know about array1[x]?
 - $\text{array1}[x] \text{ ~~*4K~~} = 64 * \text{target_set} + \text{some multiple of number of sets}$
 - $\text{array1}[x] \text{ ~~*4K~~} = 64 * 3 + n * 32K * 64$
 - So $\text{array1}[x] = 196 + n * 32K * 64$
 - So array1[x] = 196 or some large number
 -